

A vibrant, painterly illustration of the Seattle skyline. The Space Needle stands prominently on the left, its white spire reaching towards the top of the frame. To its right is a cluster of modern skyscrapers, including the Smith Tower and the Columbia Center. In the background, the majestic Mount Rainier rises in the distance, its peak partially obscured by clouds. The foreground features some greenery and lower buildings.

# Intro to TensorFlow

## Seattle, November 2016



Josh Gordon  
@random\_forests

# Agenda

**10am - 10:30am**

- Fibonacci sequence

**10:30am - 11:15am**

- Linear regression

**11:15am - 12pm**

- Deep Dream, Style Transfer, or TensorFlow for Poets

**12pm - 12:30pm**

- Lunch

**12:30-1:30pm**

- Basic and Deep MNIST

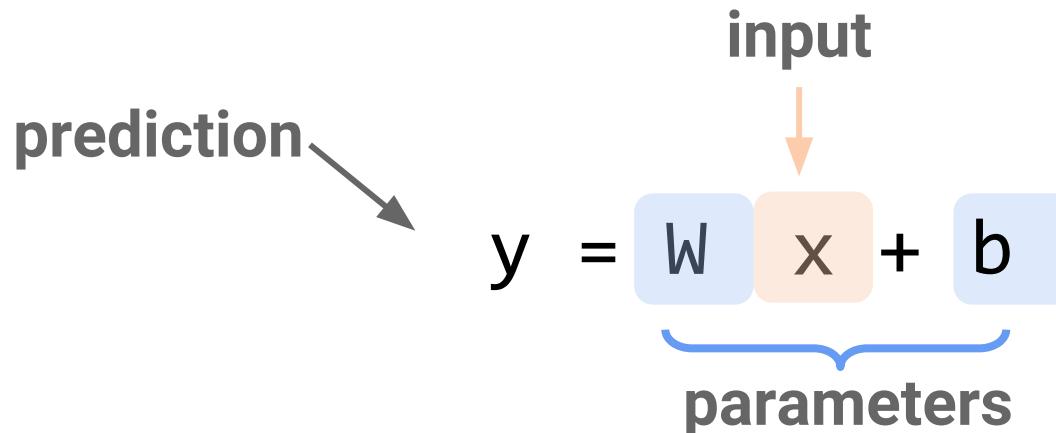
**1:30pm - 2:30pm**

- Networking

# What's Deep Learning?

# Machine Learning $\approx$ Programming + Data

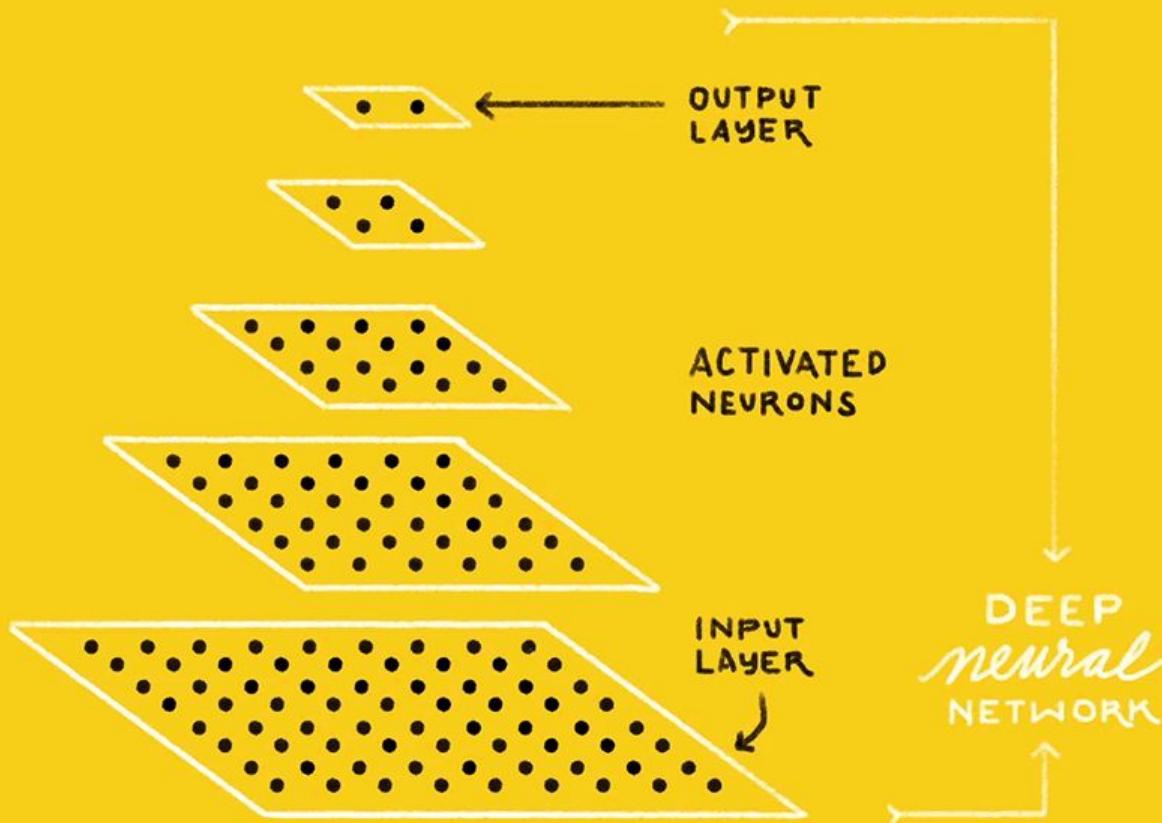
“Learn” means a data-driven procedure to find useful values for parameters.



IS THIS A  
**CAT or DOG?**



CAT   DOG



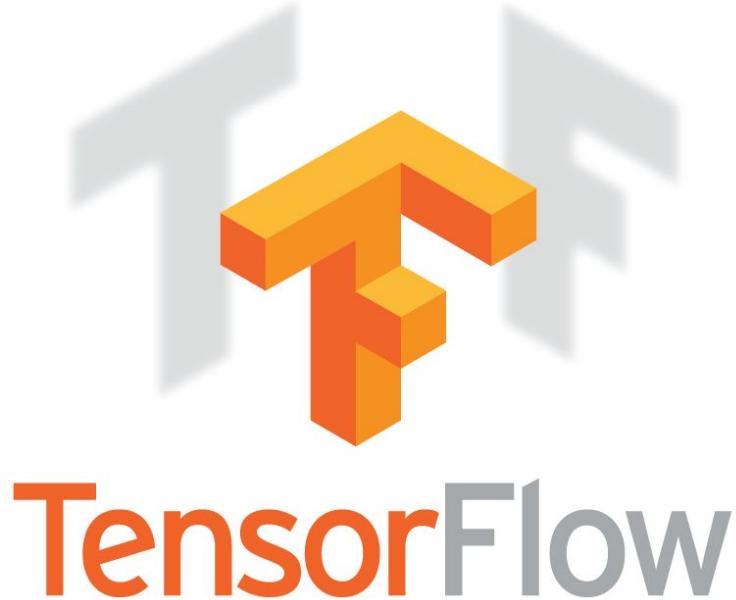
# Demo: TensorFlow Playground

[goo.gl/mXhncM](http://goo.gl/mXhncM)

# Deep Learning

Current state of the art in:

- **Image:** *classification, captioning*
- **Language:** *translation, parsing, summarization*
- **Speech:** *recognition, generation*
- **Games:** *AlphaGo, Atari*
- And much more.

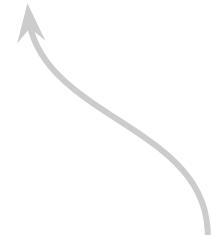


- Fast, flexible, and scalable open-source machine learning library
- One system for research and production
- Runs on CPU, GPU, TPU, and Mobile
- Apache 2.0 license

A multidimensional array.



# TensorFlow



A graph of operations.



# Open Source Models

[github.com/tensorflow/models](https://github.com/tensorflow/models)

# Inception



An Alaskan Malamute (left) and a Siberian Husky (right). Images from Wikipedia.

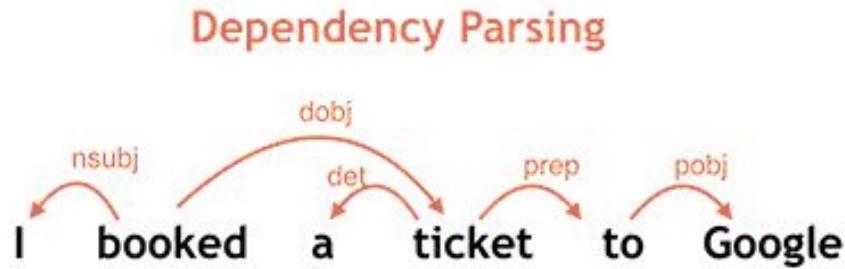
<https://research.googleblog.com/2016/08/improving-inception-and-image.html>

# Show and Tell



<https://research.googleblog.com/2016/09/show-and-tell-image-captioning-open.html>

# Parsey McParseface



<https://research.googleblog.com/2016/05/announcing-syntaxnet-worlds-most.html>

# Text Summarization

## Original text

- *Alice and Bob took the train to visit the zoo. They saw a **baby giraffe**, a **lion**, and a **flock of colorful tropical birds**.*

## Abstractive summary

- *Alice and Bob visited the zoo and saw **animals and birds**.*



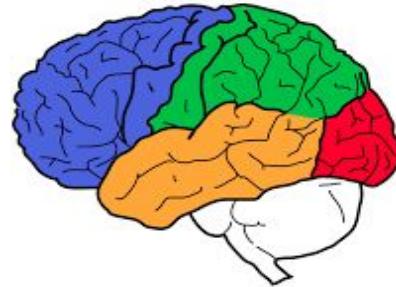
<https://github.com/lengstrom/fast-style-transfer/>



<https://www.youtube.com/watch?v=DgPaCWJL7XI>

# “Universal” Machine Learning

*Speech  
Text  
Search  
Queries  
Images  
Videos  
Labels  
Entities  
Words  
Audio  
Features*



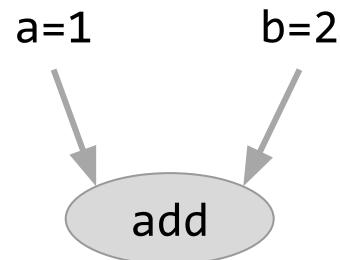
*Speech  
Text  
Search  
Queries  
Images  
Videos  
Labels  
Entities  
Words  
Audio  
Features*

# Warm up

## Graphs and Sessions

# Build a graph...

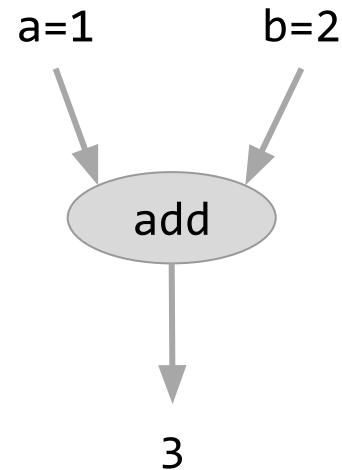
```
import tensorflow as tf  
a = tf.constant(1)  
b = tf.constant(2)  
c = tf.add(a, b)
```



# ... then run it

```
import tensorflow as tf  
a = tf.constant(1)  
b = tf.constant(2)  
c = tf.add(a, b)
```

```
session = tf.Session()  
value_of_c = session.run(c)
```



# Building graphs looks *mostly* like numpy

With special operations for deep learning

Numpy	TensorFlow
add	add
mul	mul
matmul	matmul
...	...
sum	reduce_sum
...	...
	sigmoid
	relu
	...

# What's the output?

```
import tensorflow as tf

my_var = tf.Variable(0)
add_op = tf.add(my_var,1)
assign_op = tf.assign(my_var, add_op)

session = tf.Session()
session.run(tf.initialize_all_variables())

for _ in range(3):
    print (session.run(assign_op))
```

These slides + code  
[goo.gl/nrdsxM](http://goo.gl/nrdsxM)

# Install TensorFlow

```
$ git clone https://github.com/random-forests/tensorflow-workshop  
$ cd tensorflow-workshop  
$ jupyter notebook
```

[goo.gl/nrdsxM](http://goo.gl/nrdsxM)

# Exercise #1: Fibonacci Sequence

## 1\_warm\_up.ipynb

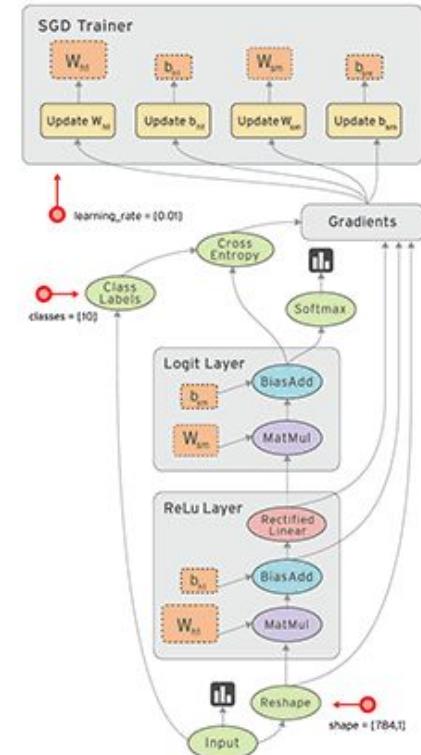
[goo.gl/nrdsxM](http://goo.gl/nrdsxM)



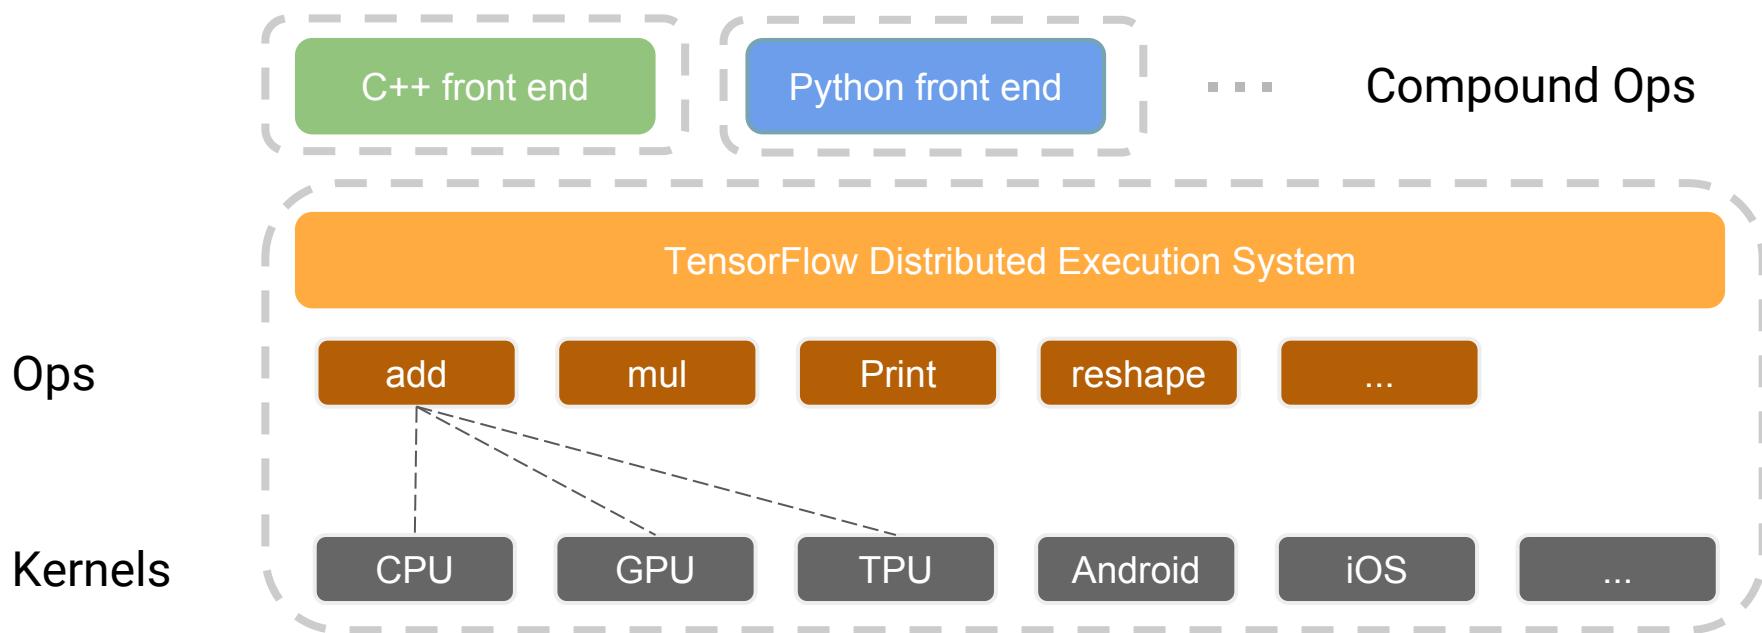
But why?

# A graph is like a blueprint

- You define the graph in high-level language (Python)
- Graph is compiled and optimized
- Graph is executed (in parts or fully) on available devices (CPU, GPU, TPU)
- Nodes represent computations
- Data (tensors) flows between them



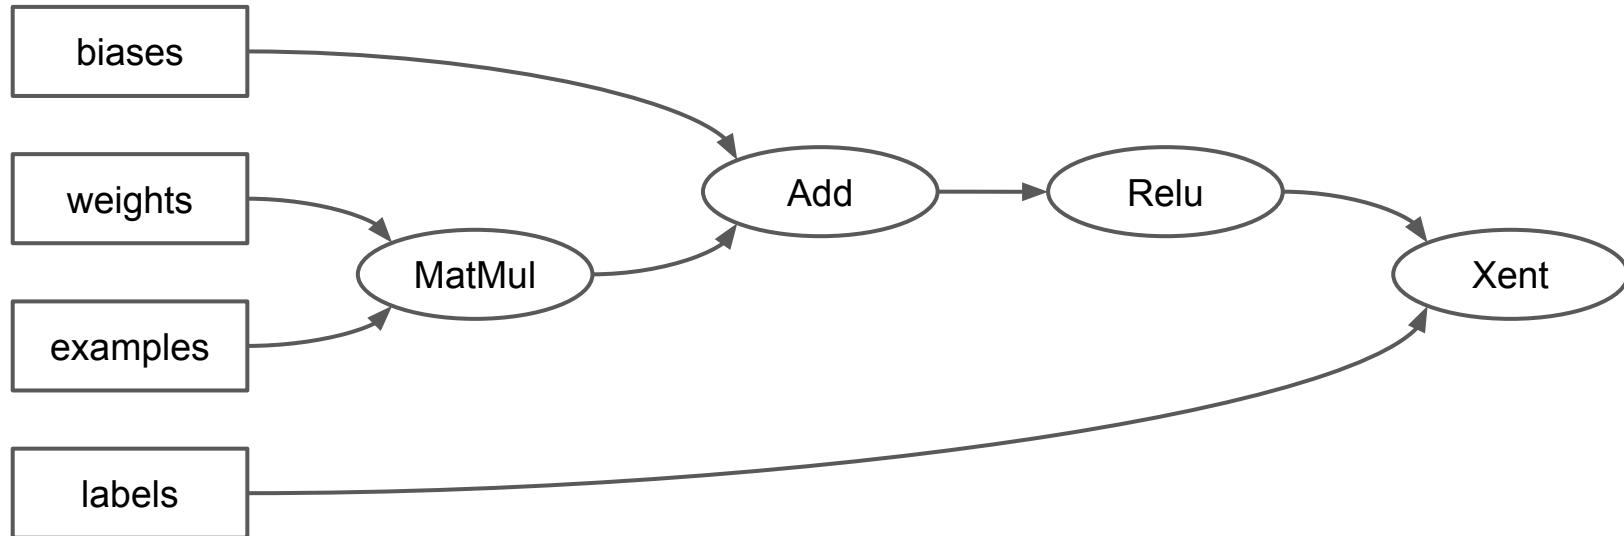
# Architecture



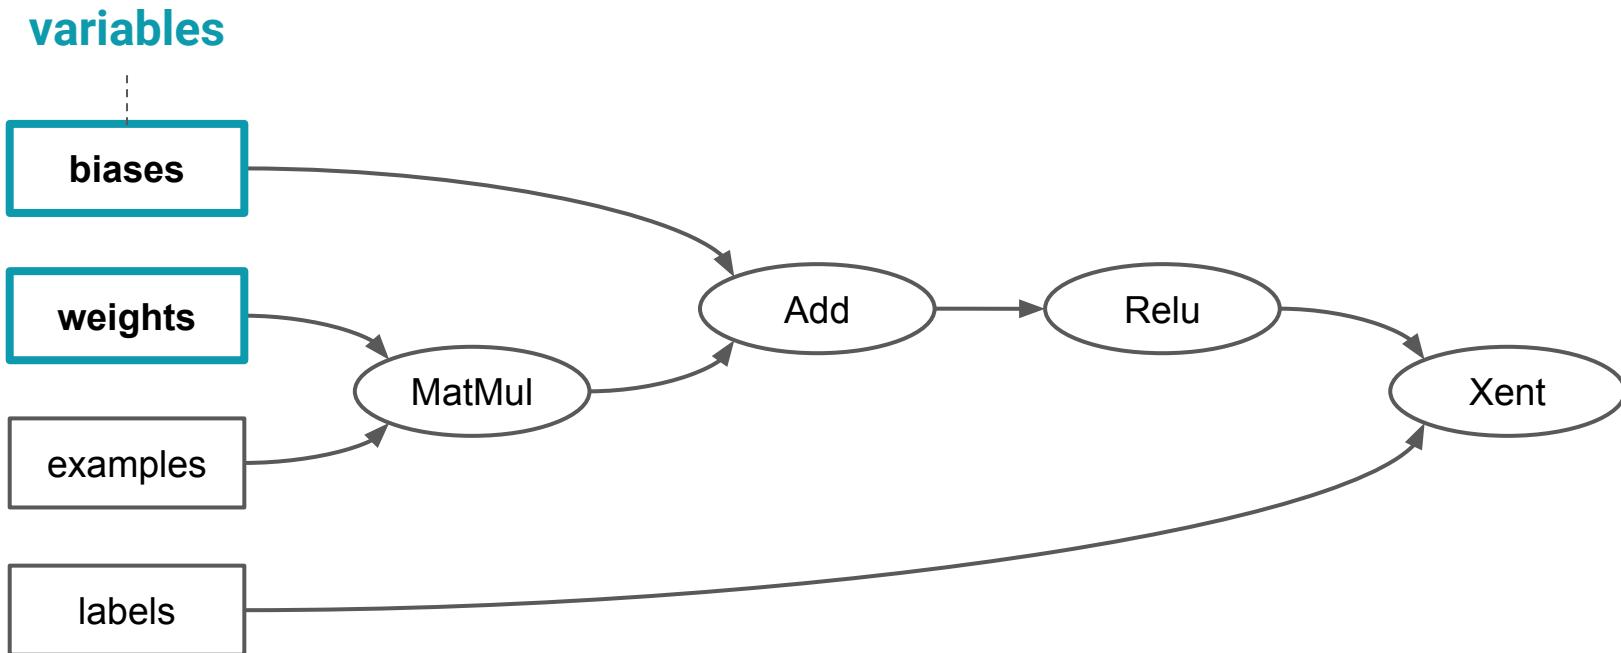
# A Year of TensorFlow

- 0.6 Python 3.3; Faster on GPUs
- 0.7 TensorFlow Serving
- 0.8 Distributed TensorFLOW
- 0.9 iOS, Raspberry Pi
- 0.10 Slim
- 0.11 CuDNN v5

**Graphs can be processed, compiled, remotely executed, assigned to devices.**

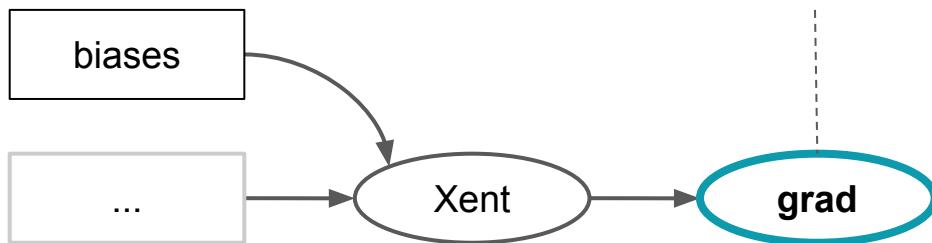


# Data flow, with state



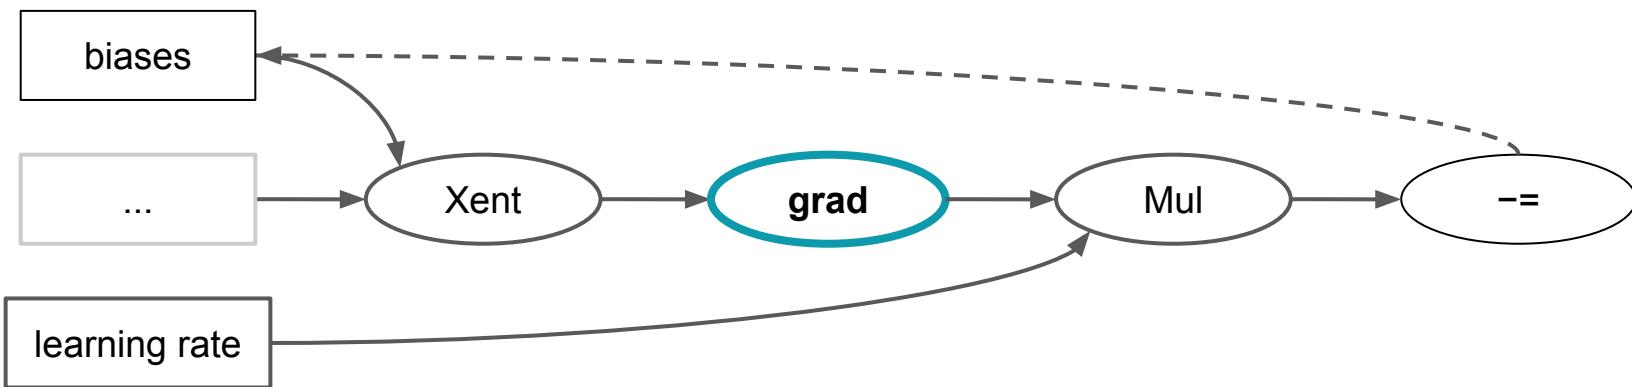
# Automatic Differentiation

Automatically add ops which  
compute gradients for variables

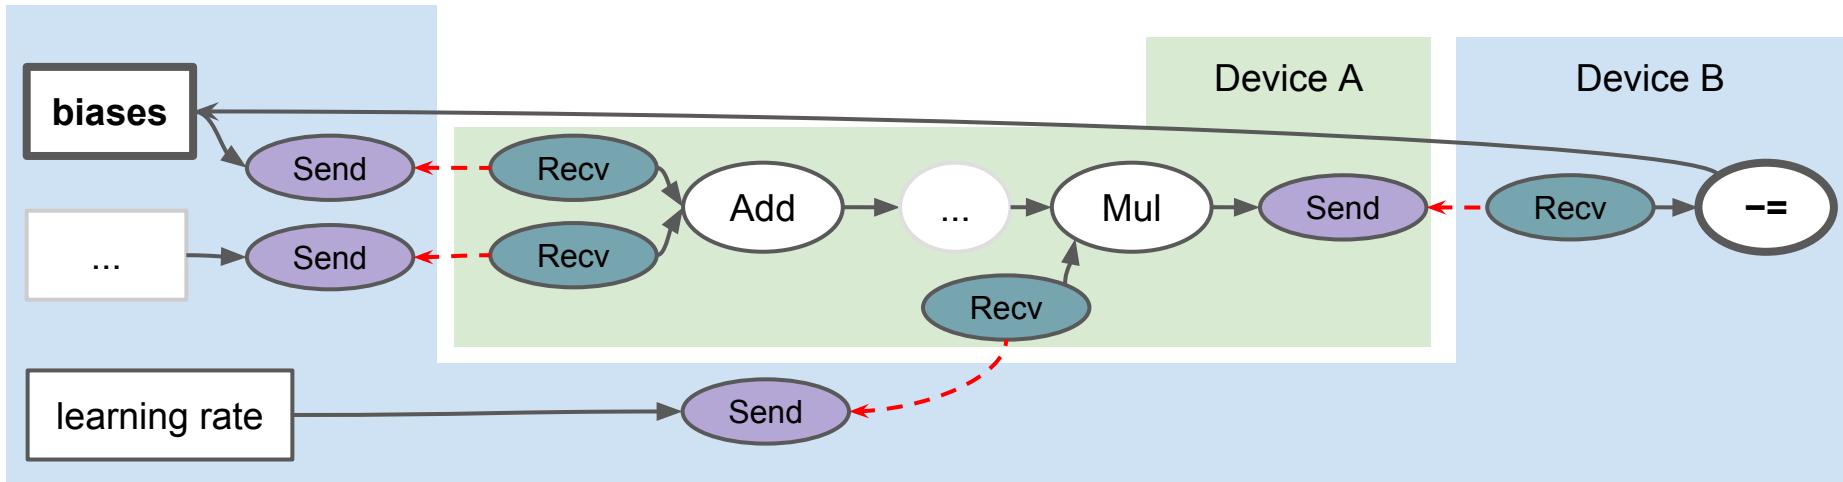


# Automatic Differentiation

Automatically add ops which compute gradients for variables



# Distributed



Send / receive nodes added automatically

Devices: Processes, Machines, CPUs, GPUs, TPUs, etc

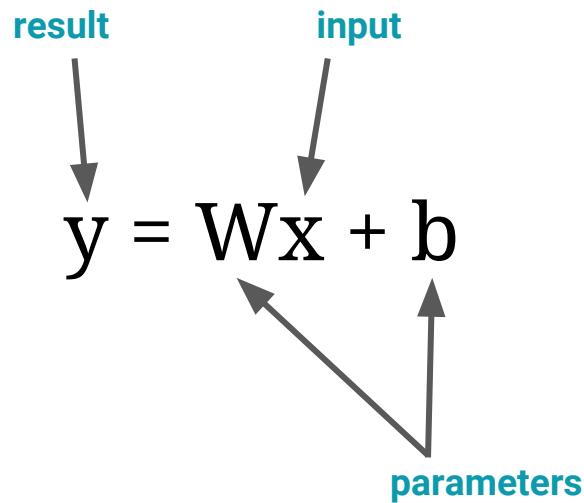
# Linear Regression

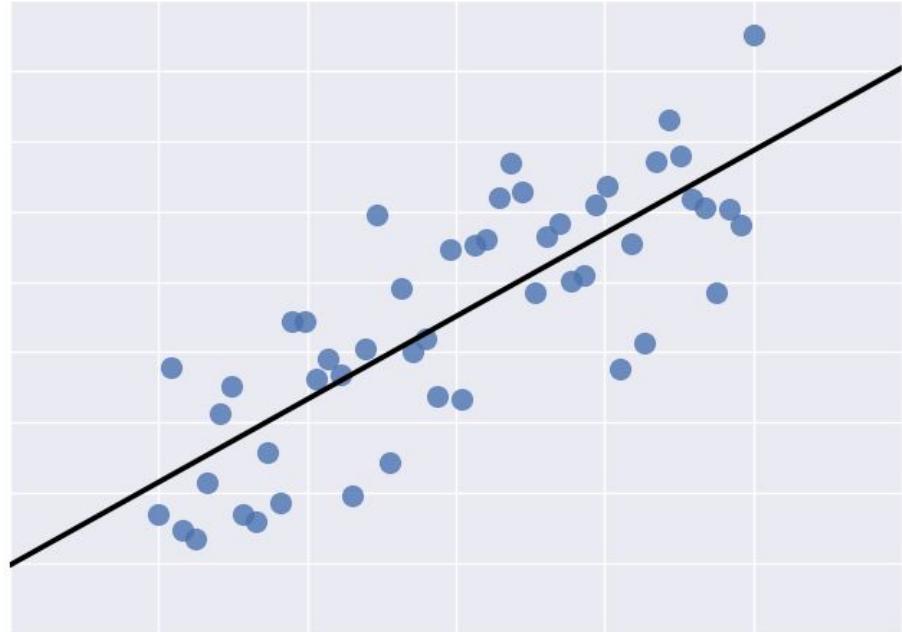
## Inference, Loss, Training

# Linear Regression

$$y = Wx + b$$

result      input  
parameters





# What are we trying to do?

**Mystery equation:**  $y = 0.1 * x + 0.3 + \text{noise}$

**Model:**  $y = W * x + b$

**Objective:** Given enough  $(x, y)$  samples, figure out the value of  $W$  and  $b$ .

# Let's code this up

1. **Inference** (“given  $x$ , this is the code to predicts  $y$ ”)
2. **Loss** (“quantify how bad our prediction was”)
3. **Optimize** (“update variables to improve the prediction”)

# $y = Wx + b$ in TensorFlow

```
import tensorflow as tf
```

# $y = Wx + b$ in TensorFlow

```
import tensorflow as tf  
  
x = tf.placeholder(shape=[None],  
                   dtype=tf.float32, name="x")
```

# $y = Wx + b$ in TensorFlow

```
import tensorflow as tf

x = tf.placeholder(shape=[None],
                    dtype=tf.float32, name="x")

W = tf.Variable(tf.random_normal([1], name="W")
```

# $y = Wx + b$ in TensorFlow

```
import tensorflow as tf

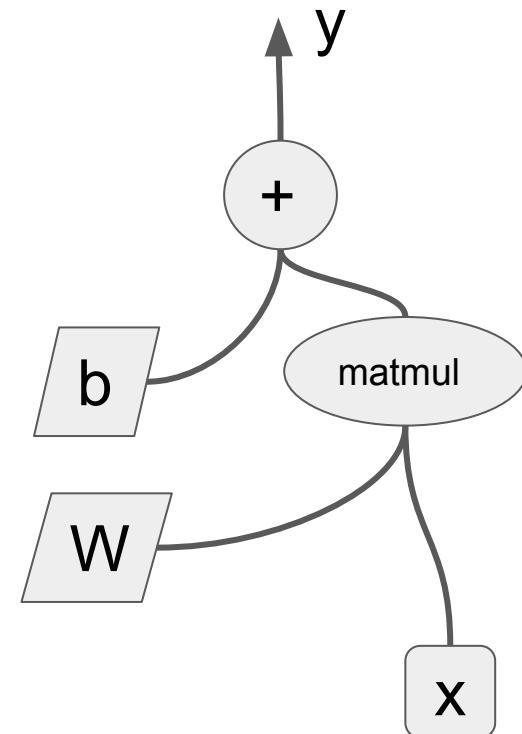
x = tf.placeholder(shape=[None],
                    dtype=tf.float32, name="x")

W = tf.Variable(tf.random_normal([1], name="W"))

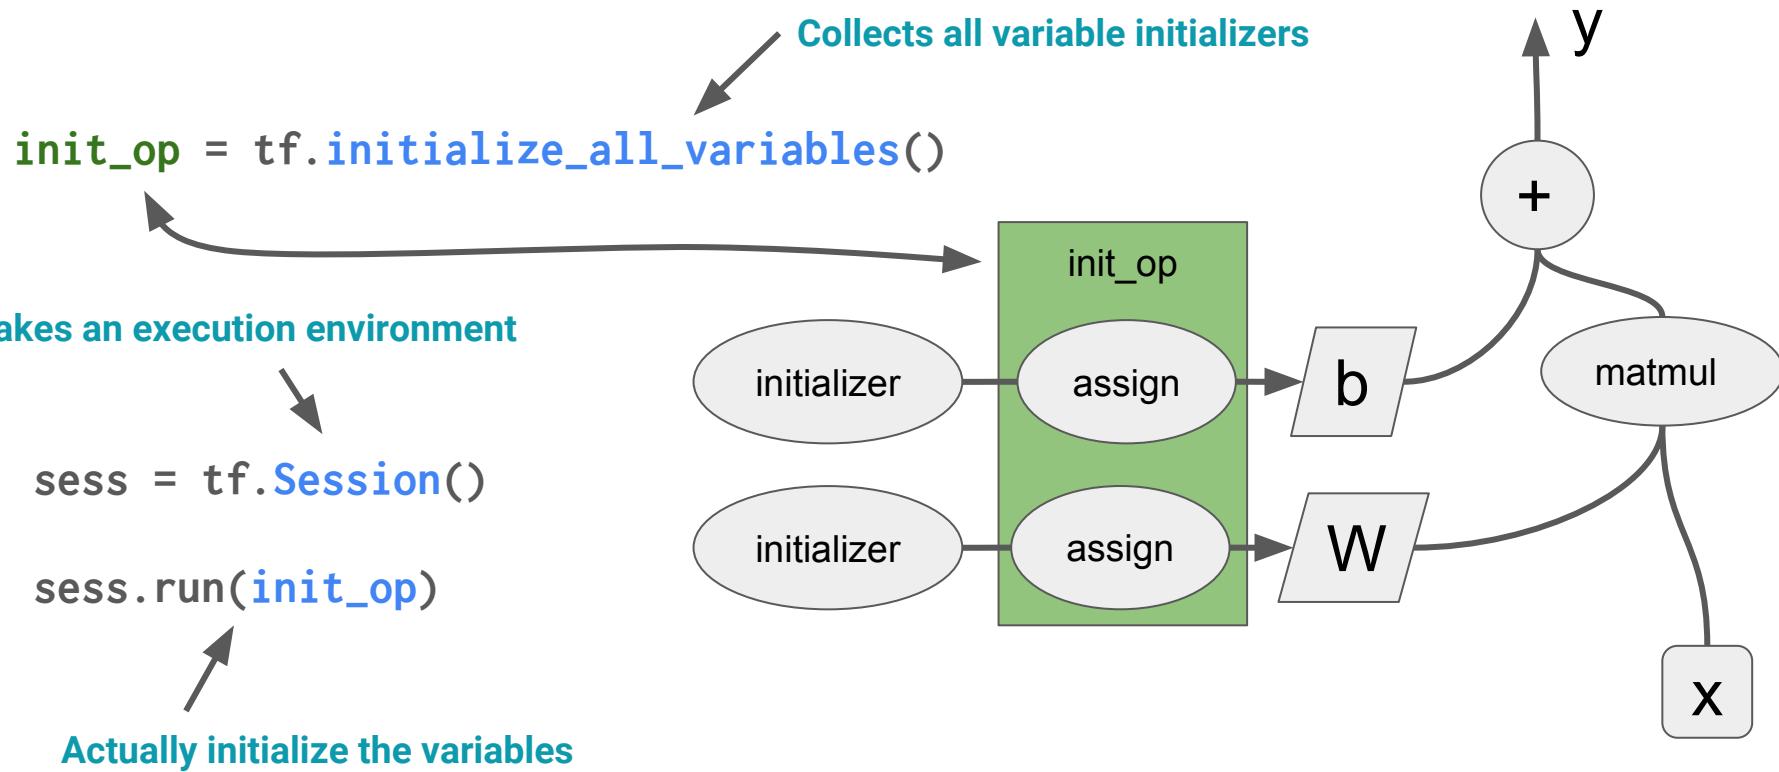
b = tf.Variable(tf.random_normal([1], name="b"))
```

# $y = Wx + b$ in TensorFlow

```
import tensorflow as tf  
  
x = tf.placeholder(shape=[None],  
                    dtype=tf.float32, name="x")  
  
W = tf.Variable(tf.random_normal([1], name="W")  
  
b = tf.Variable(tf.random_normal([1], name="b")  
  
y = W * x + b
```



# Variables Must be Initialized

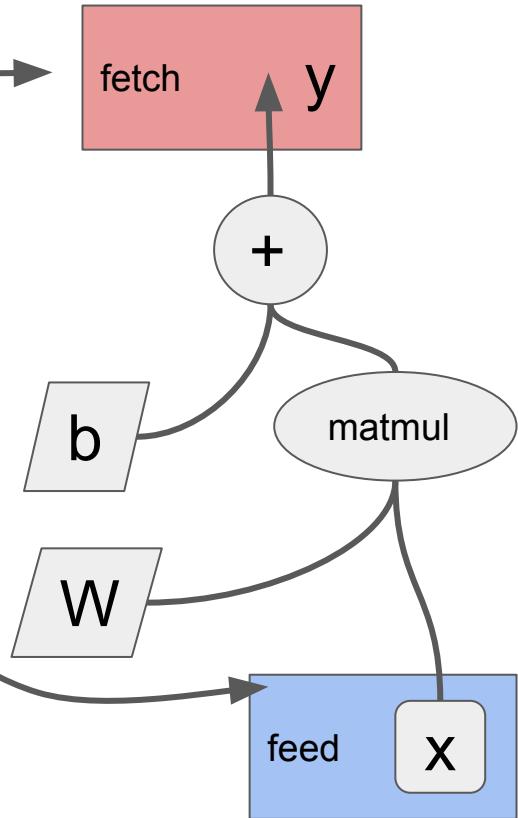


# Running the Computation

`x_in = 3`

`sess.run(y, feed_dict={x: x_in})`

- Only what's used to compute a fetch will be evaluated
- All Tensors can be fed, but all placeholders must be fed



# Inference: putting it all together

```
import tensorflow as tf
x = tf.placeholder(shape=[None],
                    dtype=tf.float32,
                    name='x')
W = tf.Variable(tf.random_normal([1], name="W"))
b = tf.Variable(tf.random_normal([1], name="b"))
y = W * x + b

with tf.Session() as sess:
    sess.run(tf.initialize_all_variables())
    print(sess.run(y, feed_dict={x: x_in}))
```

The diagram illustrates the execution flow of the provided TensorFlow code. It is divided into four main phases, each indicated by a curly brace:

- Build the graph**: This phase covers the declaration of tensors and variables. In the code, this includes the definition of `x`, `W`, and `b`.
- Prepare execution env**: This phase covers the creation of a session. In the code, this is represented by the `with tf.Session() as sess:` block.
- Initialize variables**: This phase covers the initialization of variables. In the code, this is represented by the call to `sess.run(tf.initialize_all_variables())`.
- Run the computation**: This phase covers the execution of operations. In the code, this is represented by the final `print` statement.

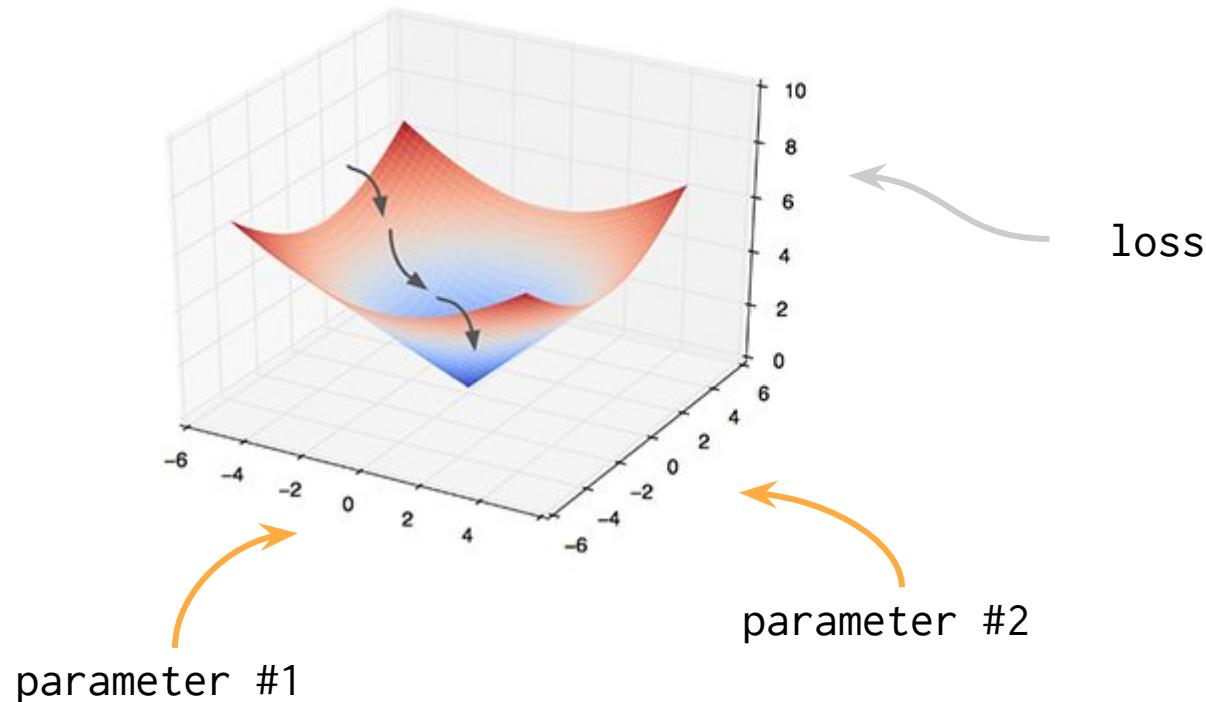
# How do we find good parameter values?

## Option 1: random guess

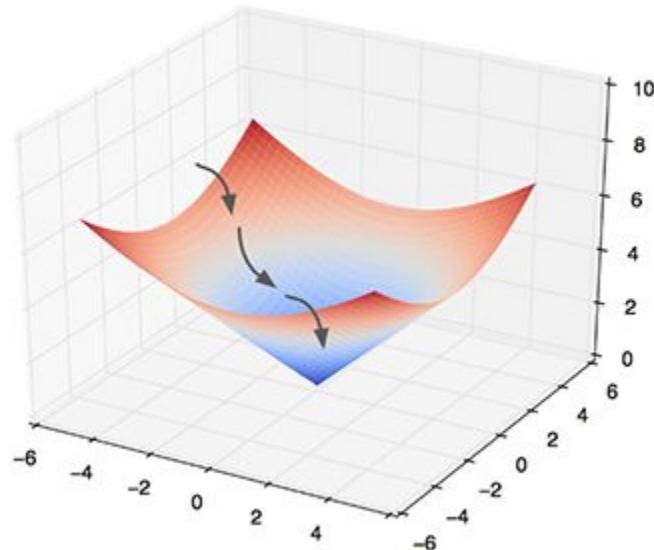
```
for _ in range(10,000)
    w,b = random(), random()
    acc = calculate_accuracy()
    if acc > best_acc:
        ...
    ...
```

## What can go wrong?

# Concepts: Loss and Gradient Descent



# Optimization



-5	-4	-3	-2	-1	0	1	2	3	4	5
5	<b>50</b>	41	34	29	26	25	26	29	34	41
4	41	<b>32</b>	<b>25</b>	20	17	16	17	20	25	32
3	34	25	18	<b>13</b>	10	9	10	13	18	25
2	29	20	13	<b>8</b>	5	4	5	8	13	20
1	26	17	10	<b>5</b>	<b>2</b>	<b>1</b>	2	5	10	17
0	25	16	9		<b>1</b>	<b>0</b>	1	4	9	16
-1	26	17	10	5	2	1	2	5	10	17
-2	29	20	13	8	5	4	5	8	13	20
-3	34	25	18	13	10	9	10	13	18	25
-4	41	32	25	20	17	16	17	20	25	32
-5	50	41	34	29	26	25	26	29	34	41

What can go wrong?

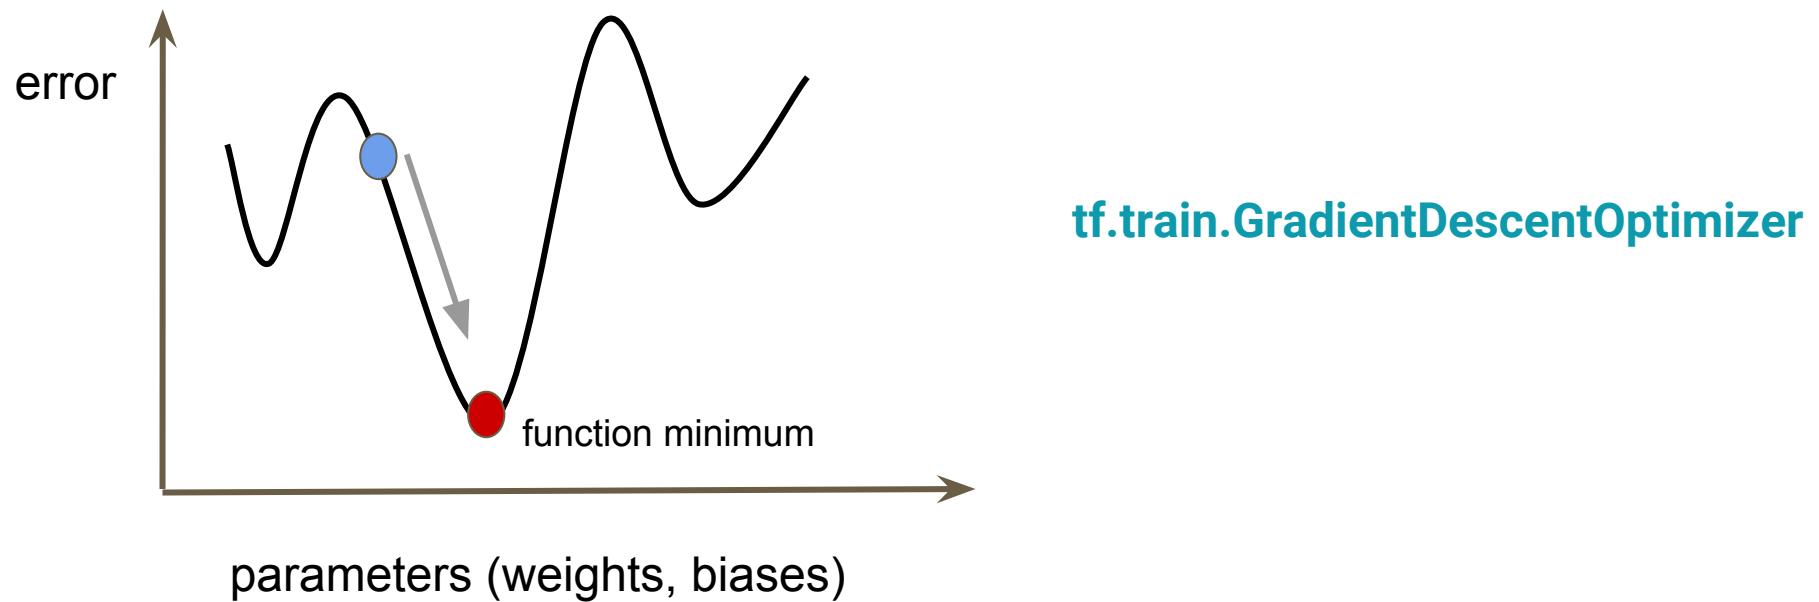
# Define a Loss

Given  $y$ ,  $y_{train}$  compute a loss, for instance:

$$loss = (y - y_{train})^2$$

```
# create an operation that calculates loss.  
loss = tf.reduce_mean(tf.square(y - y_train))
```

# Minimize loss using an optimizer



# Automatic Differentiation

Feed  $(x, y_{\text{label}})$  pairs and adjust  $W$  and  $b$  to decrease the loss.

$$W \leftarrow W - \eta ( dL/dW )$$

$$b \leftarrow b - \eta ( dL/db )$$

TensorFlow computes  
gradients automatically

```
# Create an optimizer
```

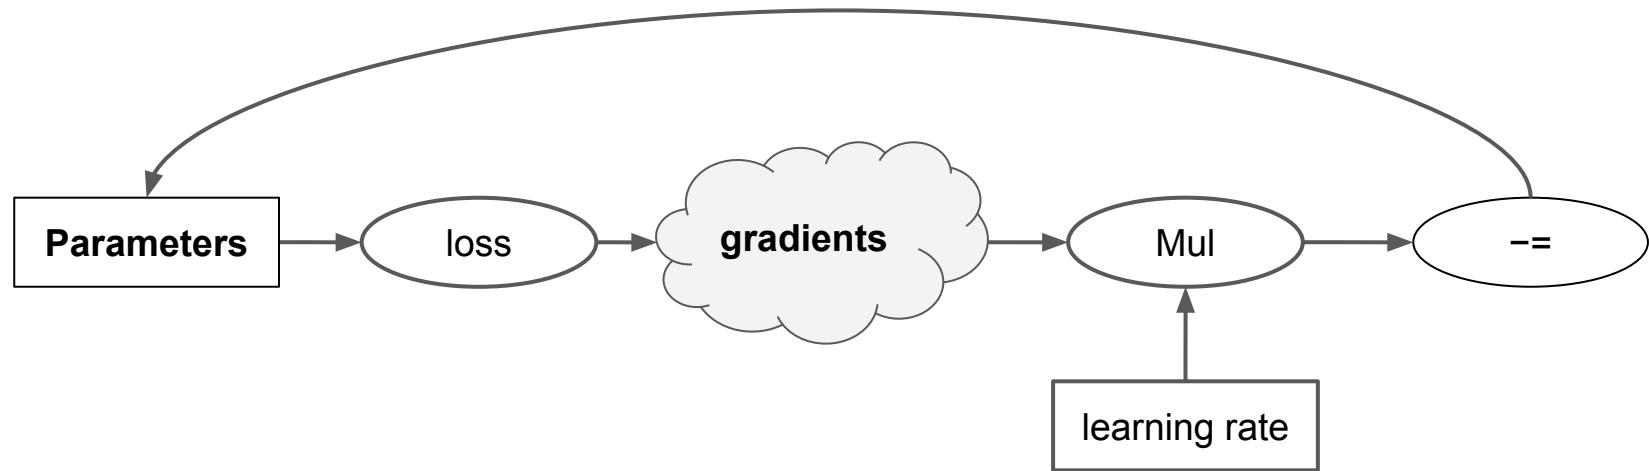
```
optimizer = tf.train.GradientDescentOptimizer(0.5)
```

```
# Create an operation that minimizes loss.
```

```
train = optimizer.minimize(loss)
```

Learning rate

# Behind the scenes



# Putting it all together

```
loss = tf.reduce_mean(tf.square(y - y_train))  
optimizer = tf.train.GradientDescentOptimizer(0.5)  
train = optimizer.minimize(loss)  
  
with tf.Session() as sess:  
  
    sess.run(tf.initialize_all_variables())  
  
    for i in range(1000):  
  
        sess.run(train, feed_dict={x: x_data[i],  
                                  y_label: y_data[i]})
```

The diagram illustrates the components of the provided TensorFlow code. It uses three horizontal braces on the right side to group parts of the code:

- A brace groups the first three lines: `loss`, `optimizer`, and `train`. It is labeled "Define a loss" (in black) above the brace.
- A brace groups the line `sess.run(tf.initialize_all_variables())`. It is labeled "Create an optimizer" (in blue) above the brace.
- A brace groups the entire loop structure starting with `for i in range(1000):`. It is labeled "Op to minimize the loss" (in green) above the brace.
- A brace groups the line `sess.run(tf.initialize_all_variables())`. It is labeled "Initialize variables" (in black) above the brace.
- A brace groups the entire loop structure starting with `for i in range(1000):`. It is labeled "Iteratively run the training op" (in orange) above the brace.

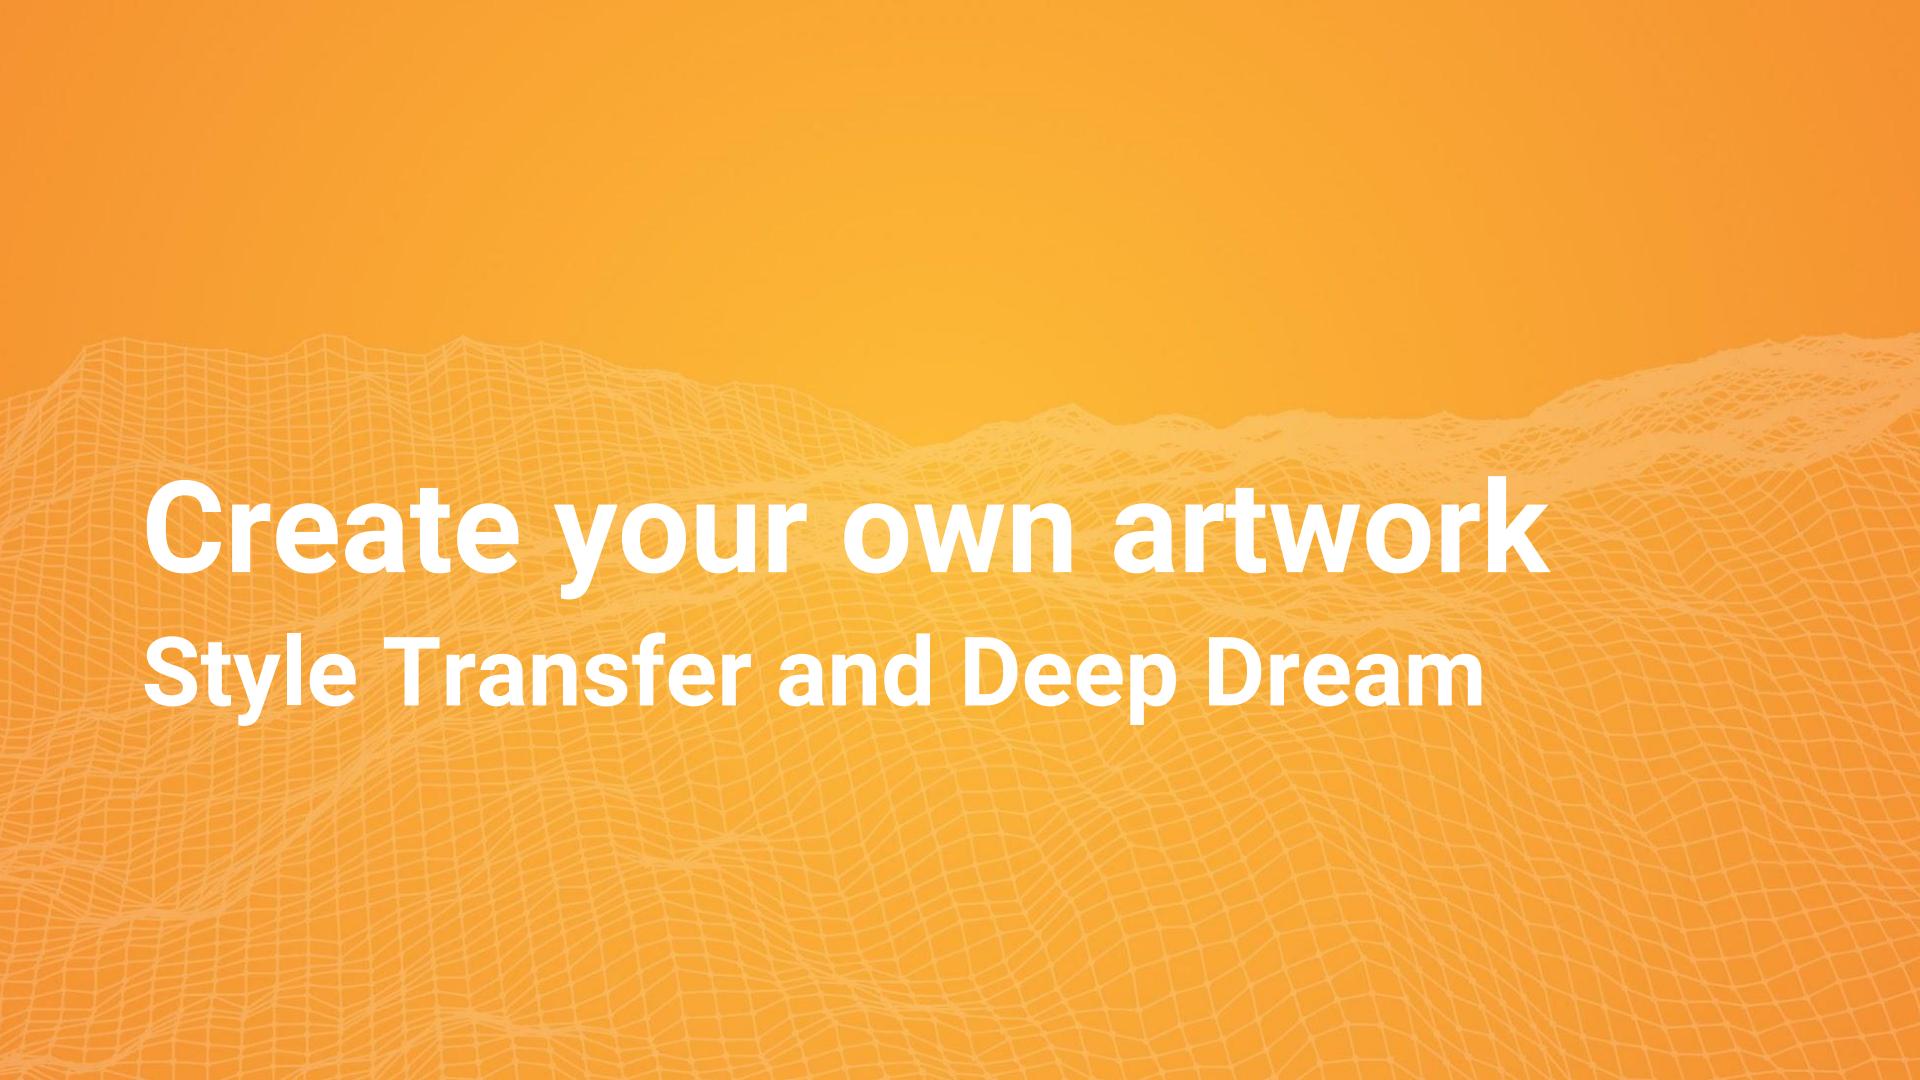
# Demo: TensorBoard

[goo.gl/nrdsxM](http://goo.gl/nrdsxM)

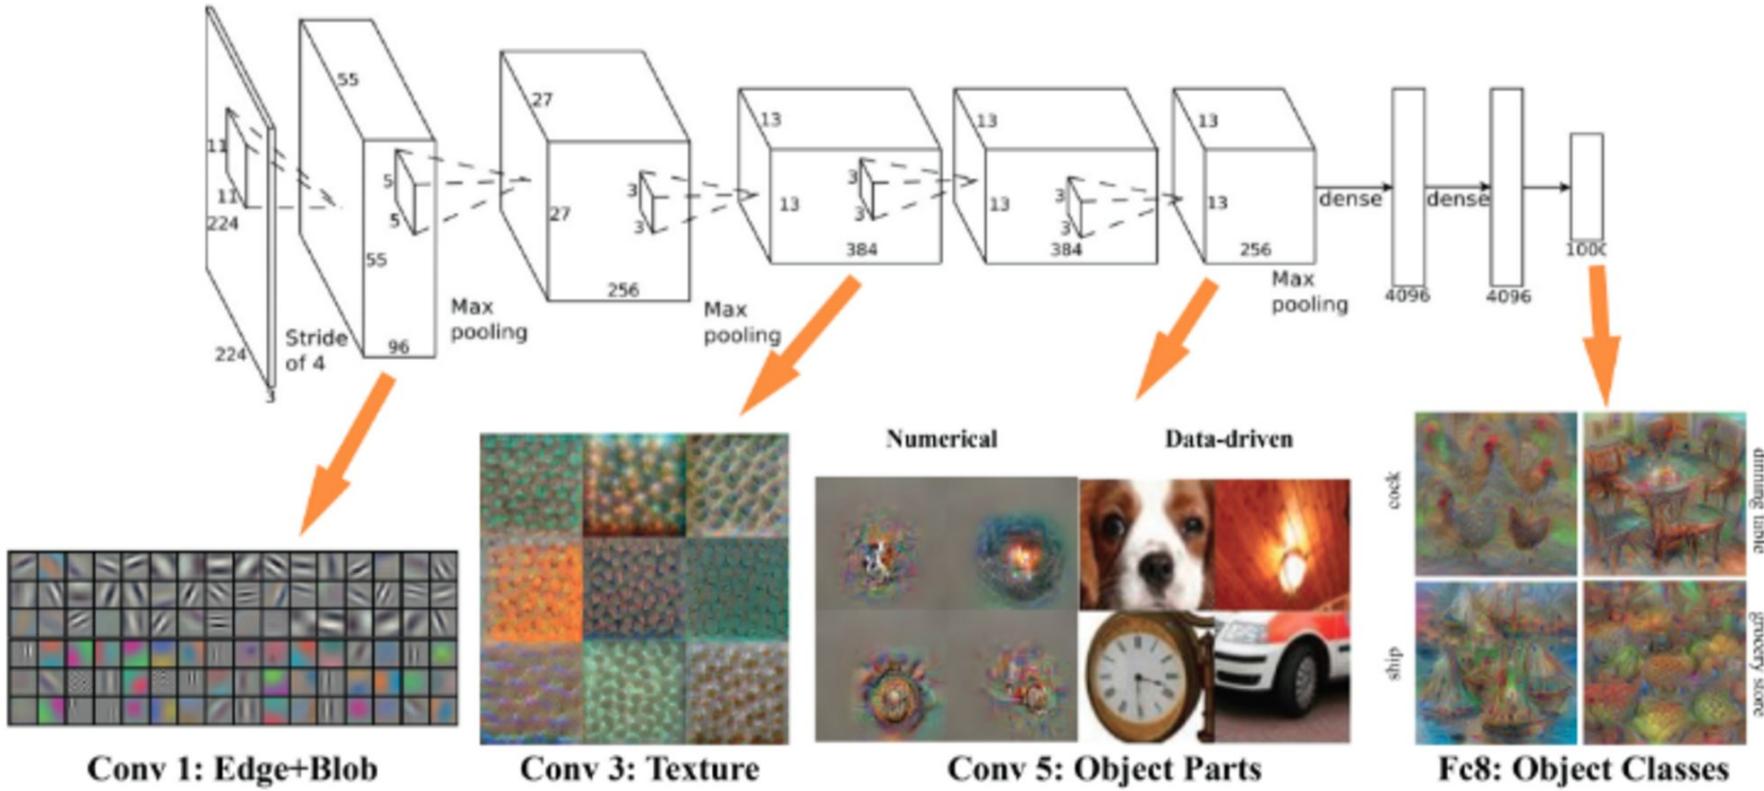
# Exercise #2: Linear Regression

## 2\_linear\_regression.ipynb

[goo.gl/nrdsxM](http://goo.gl/nrdsxM)



Create your own artwork  
Style Transfer and Deep Dream



From: [A Matlab Plugin to Visualize Neurons from Deep Models](#), Donglai Wei et. al.

# Style Transfer





Images from <https://github.com/lengstrom/fast-style-transfer/>

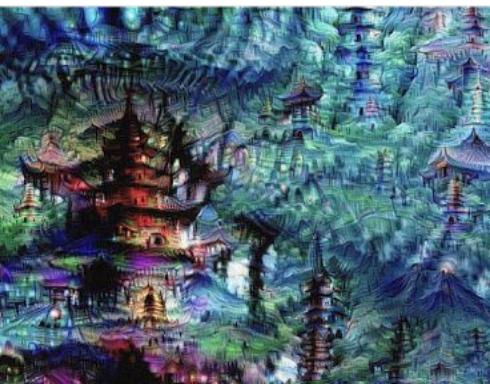
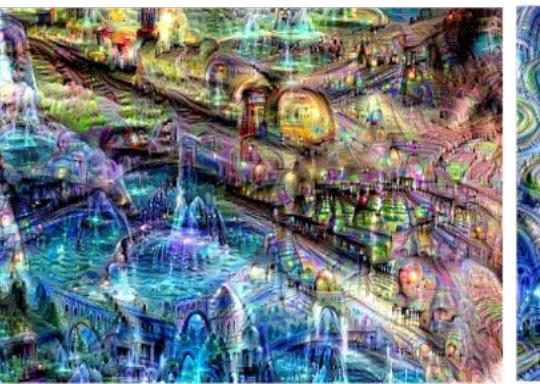
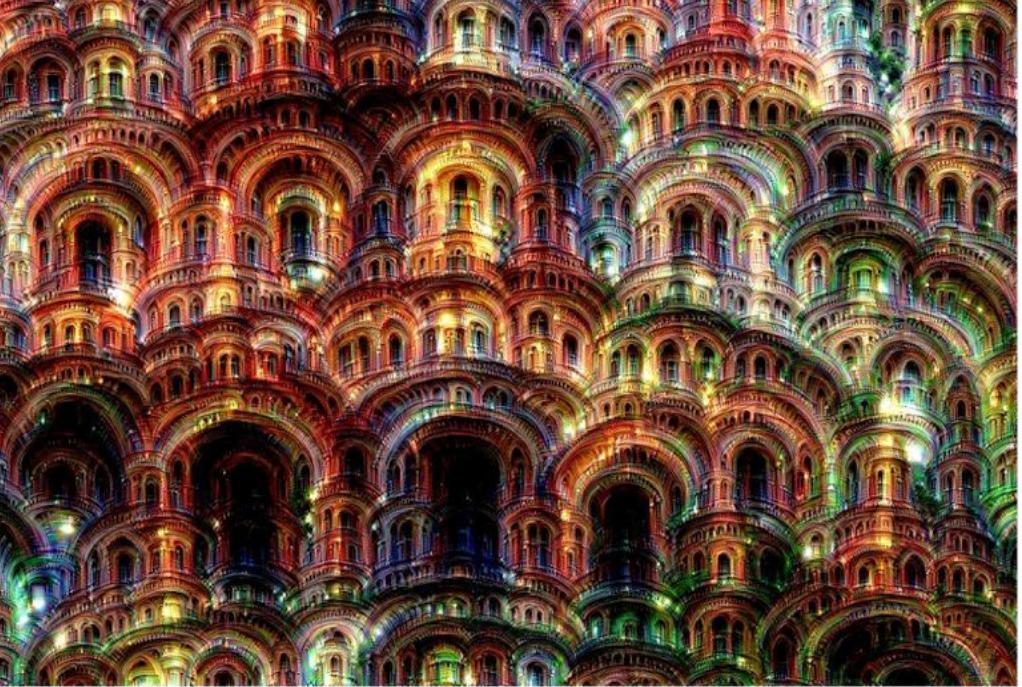




From <https://github.com/cysmith/neural-style-tf>

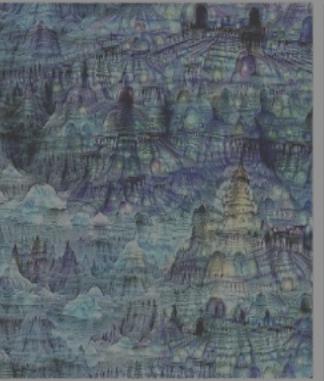
**Style Transfer**  
[goo.gl/M5hiYY](http://goo.gl/M5hiYY)

**Fast Style Transfer**  
[goo.gl/owF2z9](http://goo.gl/owF2z9)





[goo.gl/1kBXyO](http://goo.gl/1kBXyO)



# References

Supercharging Style Transfer

<https://research.googleblog.com/2016/10/supercharging-style-transfer.html>

Inceptionism: Going Deeper into Neural Networks

<https://research.googleblog.com/2015/06/inceptionism-going-deeper-into-neural.html>

# TensorFlow for Poets

## Train your own image classifier

[goo.gl/xGsB9d](http://goo.gl/xGsB9d)

# Remember Inception?

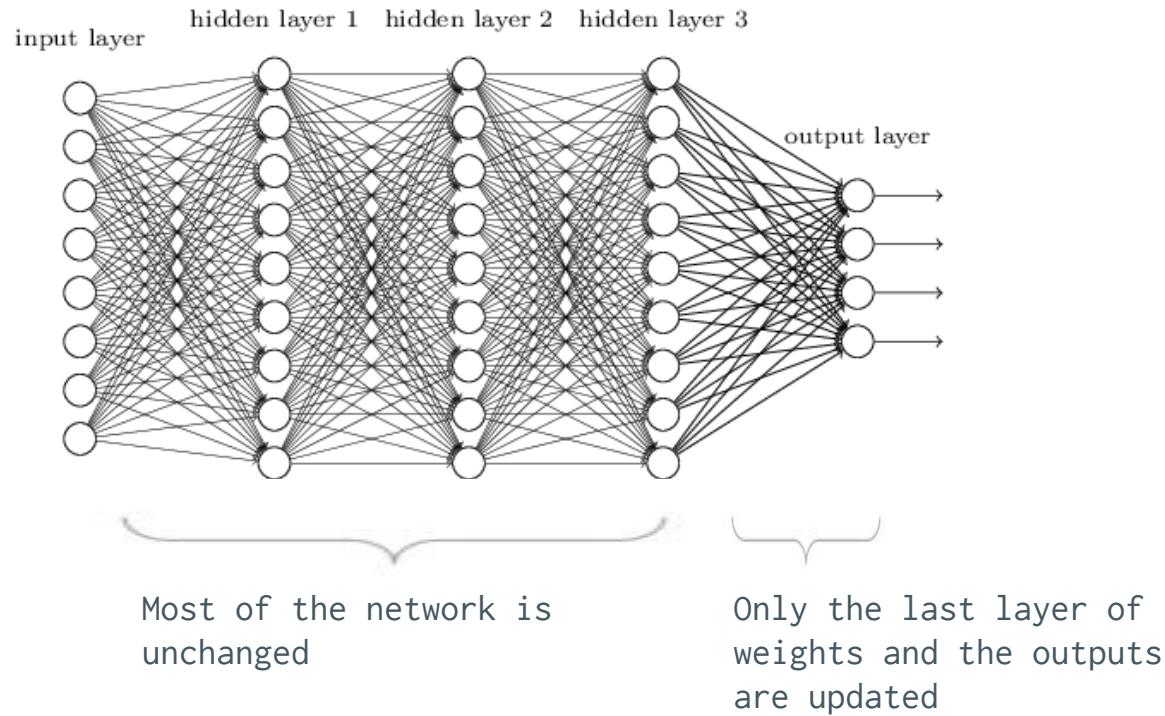
- Load the graph and learned weights; use the model
- 1,000 categories: pandas, cats, bears
- Performs “about” as well as people on this task



Images via [Wikipedia](#)

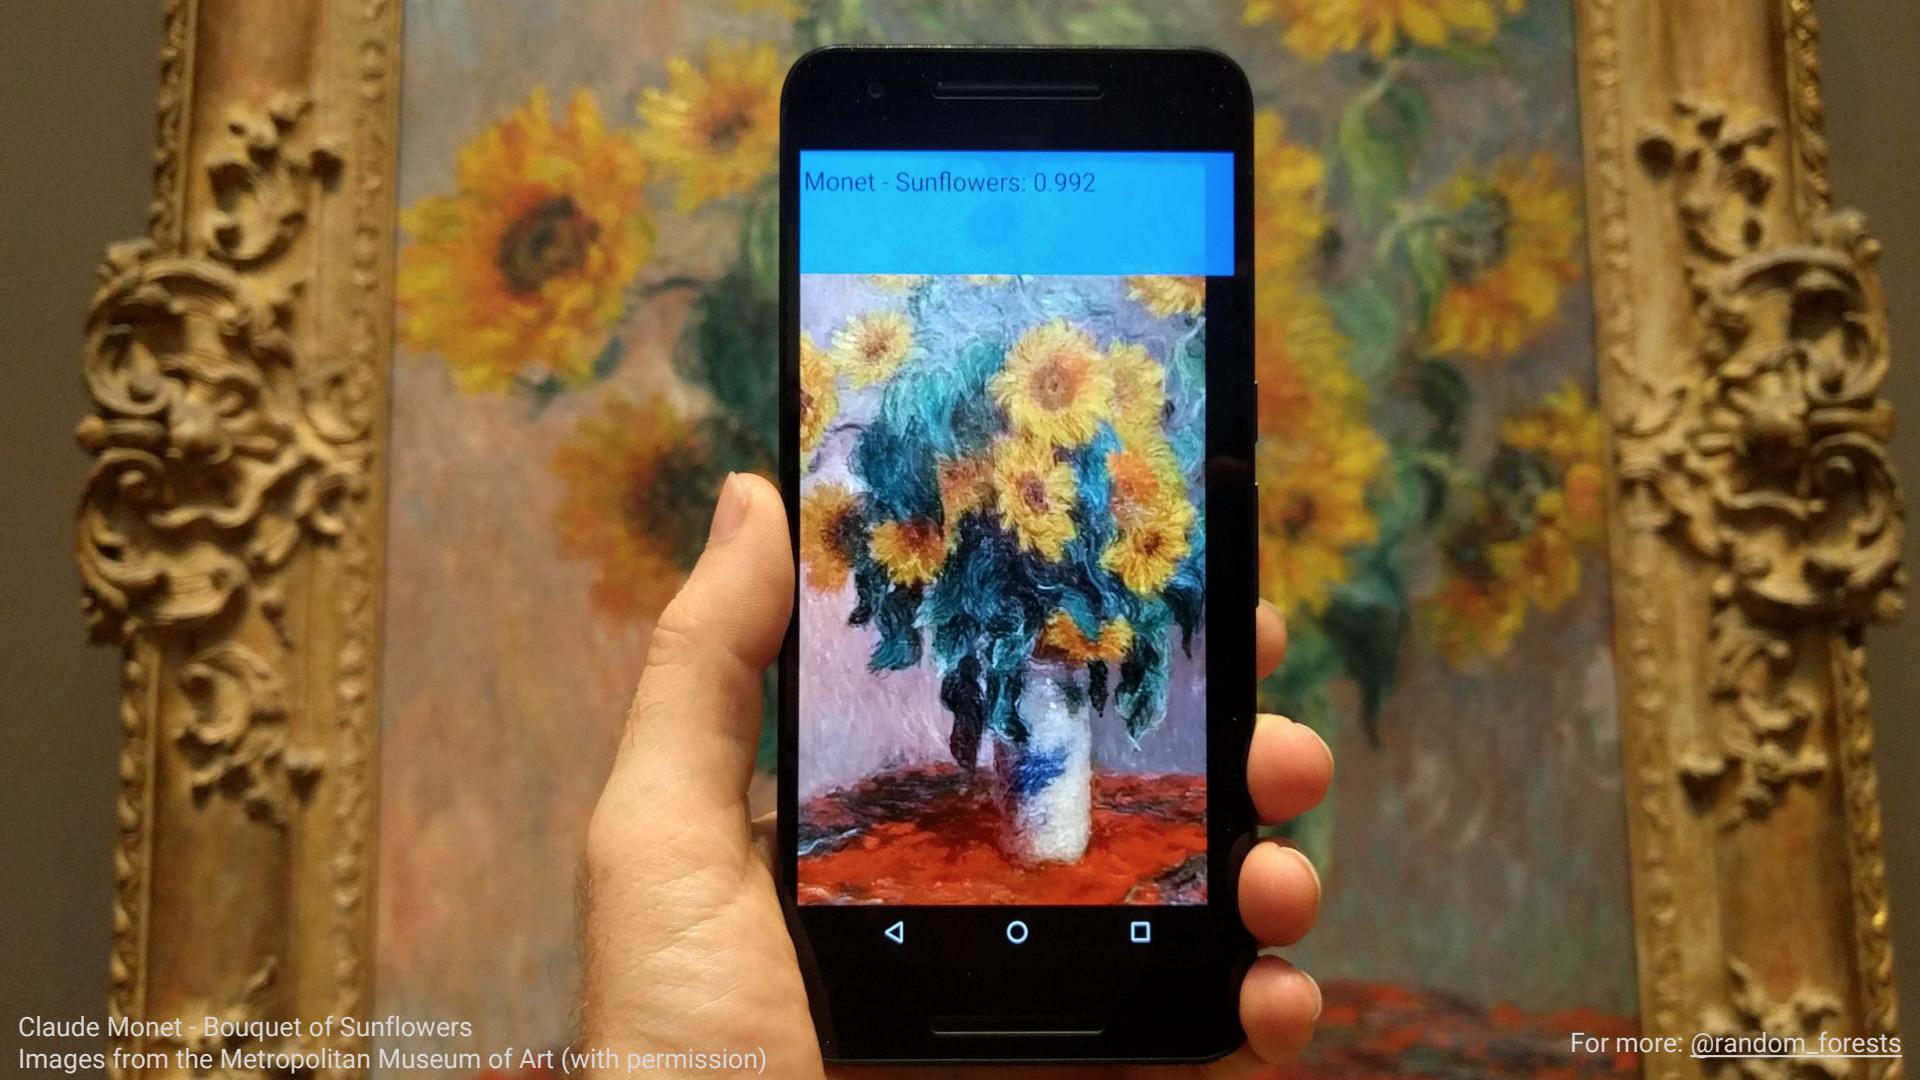
[www.tensorflow.org/tutorials/image\\_recognition/index.html](http://www.tensorflow.org/tutorials/image_recognition/index.html)

# Transfer Learning



# Advantages

- Minutes vs weeks
- Handful of new training data
- Reuse learned “features”
- Two ways to do this
  - a. Directly in TensorFlow
  - b. TensorFlow for Poets



Claude Monet - Bouquet of Sunflowers  
Images from the Metropolitan Museum of Art (with permission)

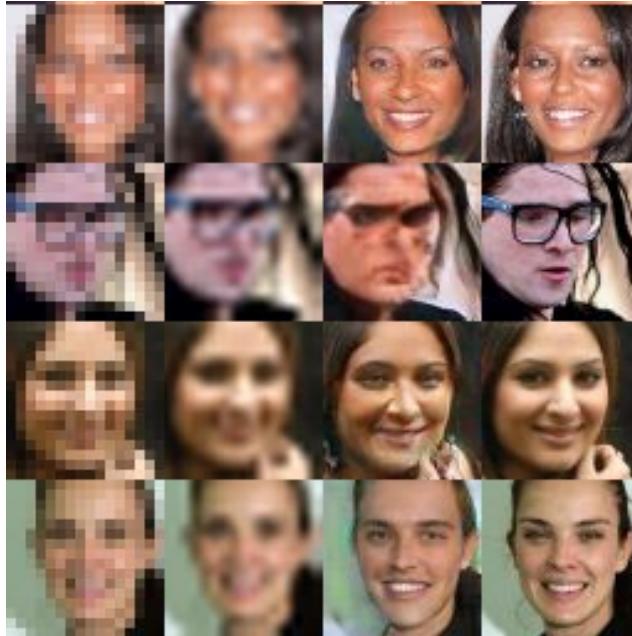
For more: [@random\\_forests](https://github.com/Random-Forests)

## **Ex #3: Divide into groups, try your favorite.**

- 1) TensorFlow for Poets**
- 2) Deep Dream**
- 3) Style Transfer**

**[goo.gl/nrdsxM](http://goo.gl/nrdsxM)**

# Also: Super-resolution



<https://github.com/david-gpu/srez>

# Basic and Deep MNIST

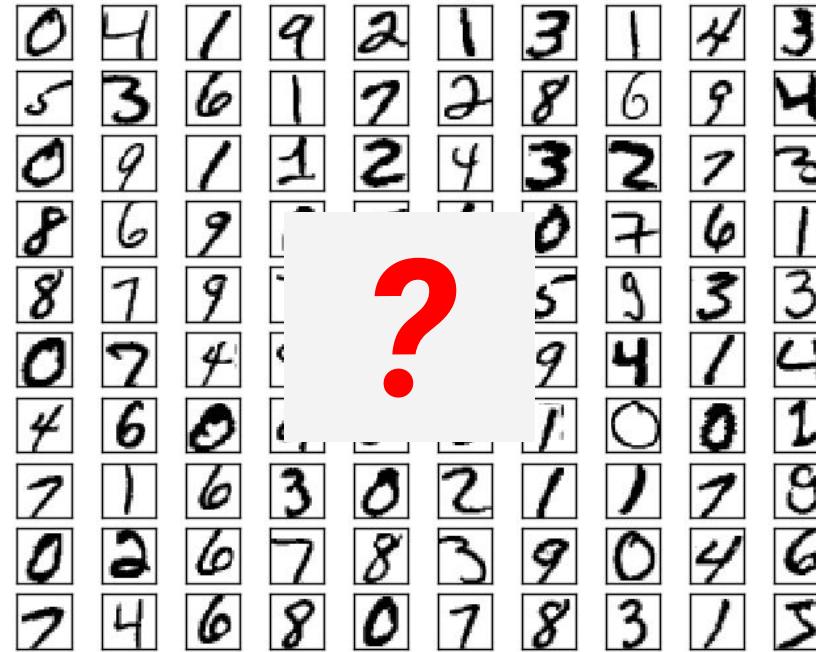
# Working with Neural Networks

1. Use a pre-trained model
2. Use a higher-level API
3. Define and train your network directly in TensorFlow

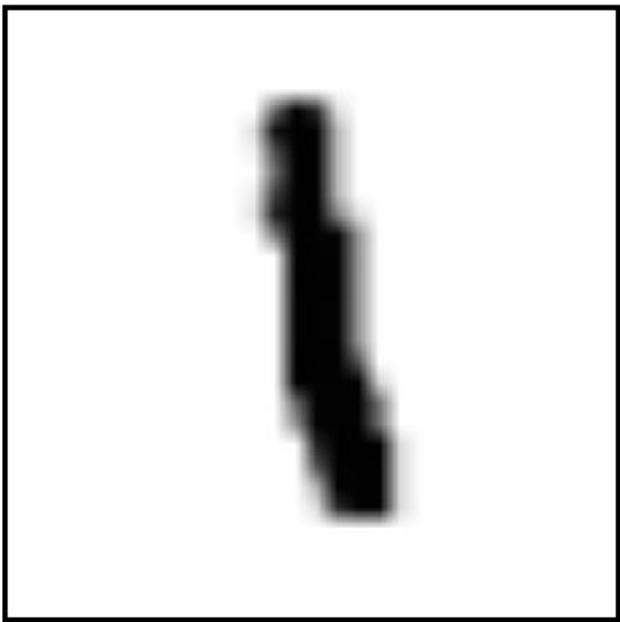


More control,  
more effort

# Hello “computer vision” world



# We see



# Computer “sees”

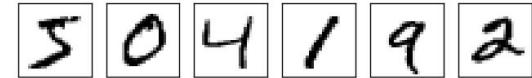
2

# Dataset

Original: 65,000

1 8 2 0 2 9 9 5 5 1 5 6 0 3 4 4 6 5 4 6 5  
9 2 5 0 1 1 1 0 9 0 3 1 6 4 2 3 6 1 1 1 3  
2 8 4 1 7 3 3 8 8 7 9 2 2 4 1 5 9 8 7 2 3  
1 8 1 8 0 3 0 1 9 9 4 1 8 2 1 2 9 7 5 9 2  
0 2 7 4 3 3 0 0 3 1 9 6 5 0 5 1 7 9 3 0 4  
0 5 1 3 1 5 5 6 1 8 5 1 1 4 4 6 2 2 5 0 6  
6 2 1 9 2 8 6 1 9 5 2 5 4 4 2 8 3 8 2 4 5  
9 1 4 8 1 8 4 5 9 9 8 3 7 6 0 0 3 0 2 6 1  
7 5 8 9 6 1 8 4 1 2 5 9 1 9 7 5 4 0 8 9 9  
6 5 7 7 2 2 6 3 2 6 5 4 8 9 7 1 3 0 3 8 3  
3 2 7 7 0 8 7 4 4 7 9 6 9 0 9 8 0 4 6 0 6  
4 3 8 0 9 6 3 8 0 9 9 6 8 6 8 5 7 8 6 0 2

Images



Labels

5 0 4 1 9 2

# Train / test split

Original: 65,000

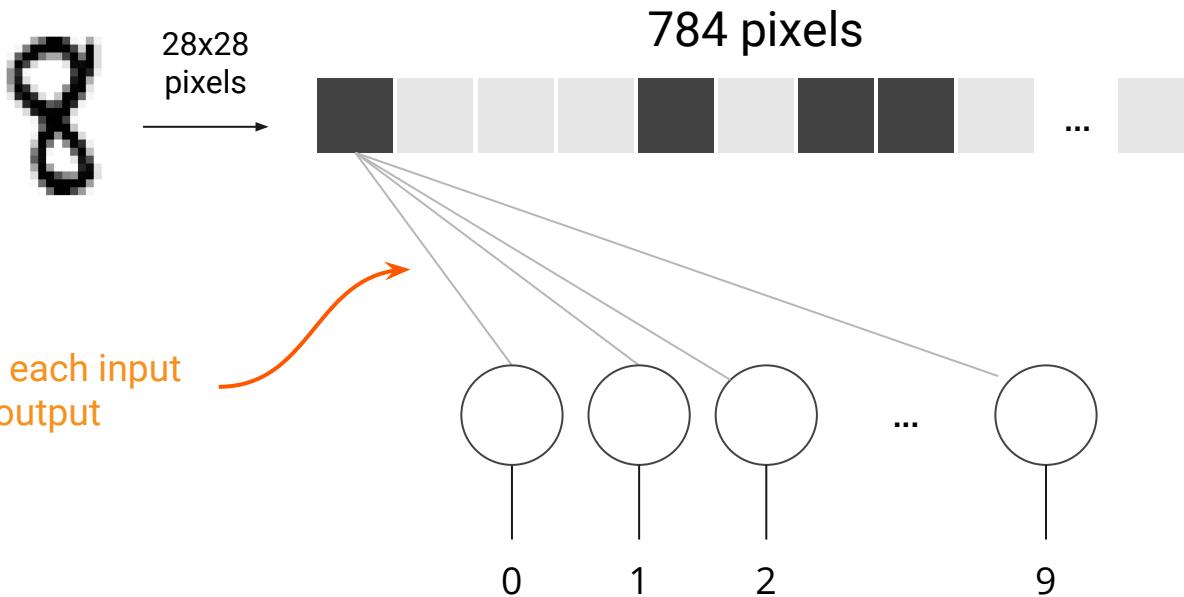
```
1 8 2 0 2 9 9 5 5 1 5 6 0 3 4 4 6 5 4 6 5  
9 2 5 0 1 1 1 0 9 0 3 1 6 4 2 3 6 1 1 1 3  
2 8 4 1 7 3 3 8 8 7 9 2 2 4 1 5 9 8 7 2 3  
1 8 1 8 0 3 0 1 9 9 4 1 8 2 1 2 9 7 5 9 2  
0 2 7 4 3 3 0 0 3 1 9 6 5 3 5 1 7 9 3 0 4  
0 5 1 3 1 5 5 6 1 8 5 1 1 4 4 6 2 2 5 0 6  
6 2 1 9 2 8 6 1 9 5 2 5 4 4 2 8 3 8 2 4 5  
9 1 4 8 1 8 4 5 9 9 8 3 7 6 0 0 3 0 2 6 4  
7 5 8 9 6 1 8 4 1 2 5 9 1 9 7 5 4 0 8 9 9  
6 5 7 1 2 6 3 2 6 5 4 8 9 7 1 3 0 3 8 3  
3 2 7 7 0 8 7 4 4 7 9 6 9 0 9 8 0 4 6 0 6  
4 3 8 0 9 6 3 8 0 9 9 6 8 6 8 5 7 8 6 0 2
```

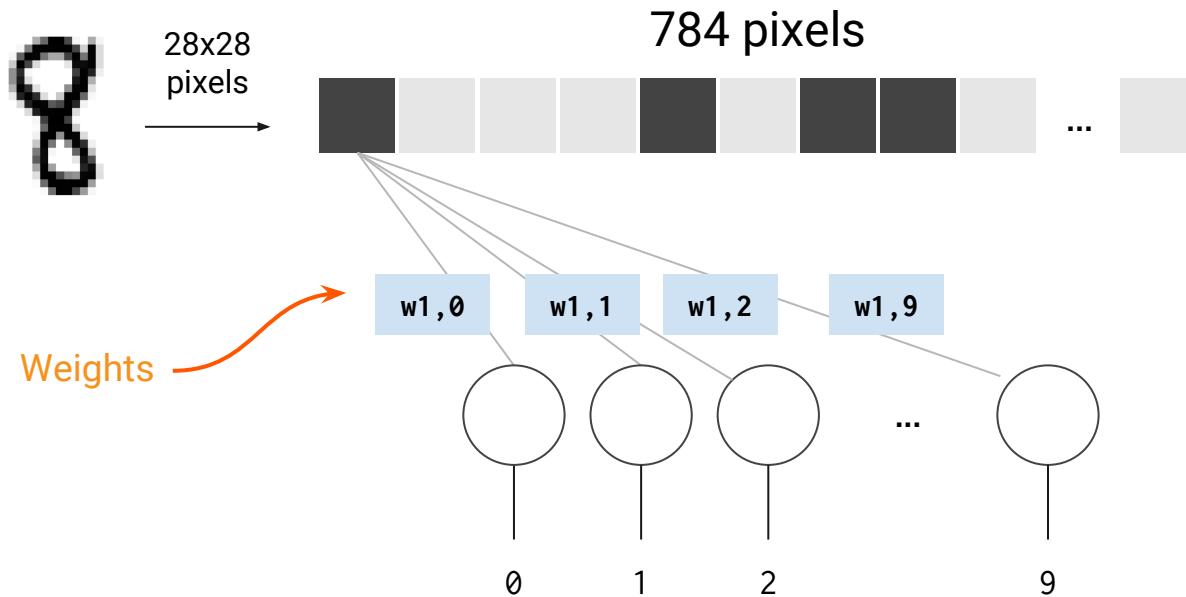
Train: 55,000

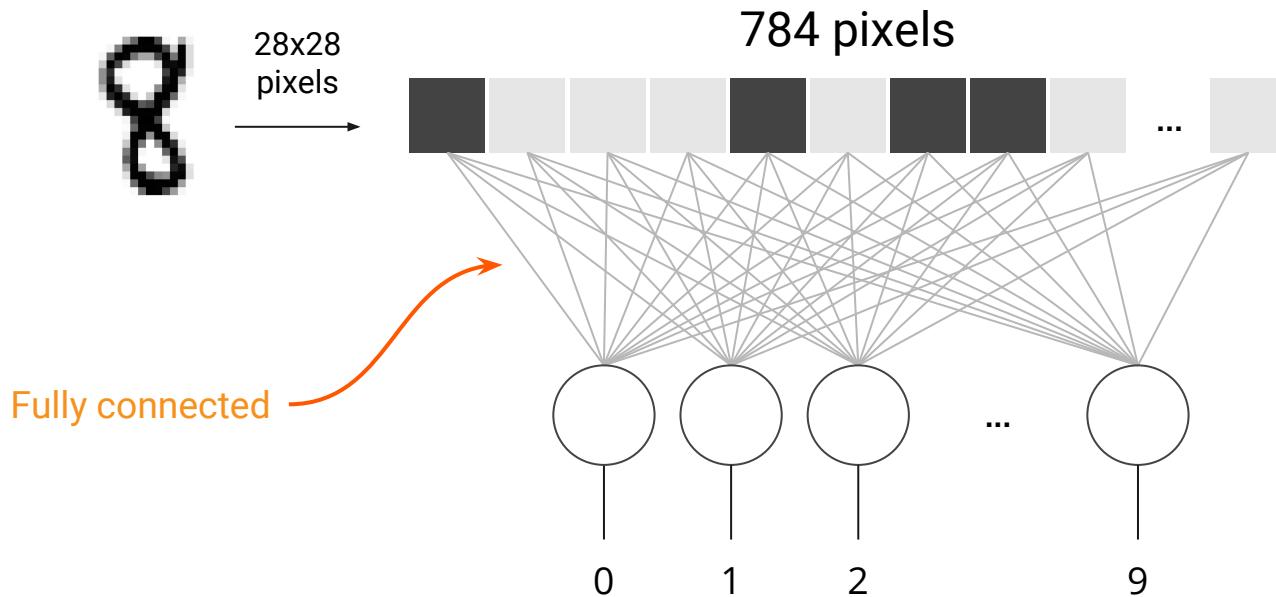
```
1 8 2 0 2 9 9 5 5 1 5 6 0 3 4 4  
9 2 5 0 1 1 1 0 9 0 3 1 6 4 2 3  
2 8 4 1 7 3 3 8 8 7 9 2 2 4 1 5  
1 8 1 8 0 3 0 1 9 9 4 1 8 2 1 2  
0 2 7 4 3 3 0 0 3 1 9 6 5 3 5 1  
0 5 1 3 1 5 5 6 1 8 5 1 1 4 4 6  
6 2 1 9 2 8 6 1 9 5 2 5 4 4 2 8  
9 1 4 8 1 8 4 5 9 9 8 3 7 6 0 0  
7 5 8 9 6 1 8 4 1 2 5 9 1 9 7 5  
6 5 7 1 2 6 3 2 6 5 4 8 9 7 1  
3 2 7 7 0 8 7 4 4 7 9 6 9 0 9 8  
4 3 8 0 9 6 3 8 0 9 9 6 8 6 8 5
```

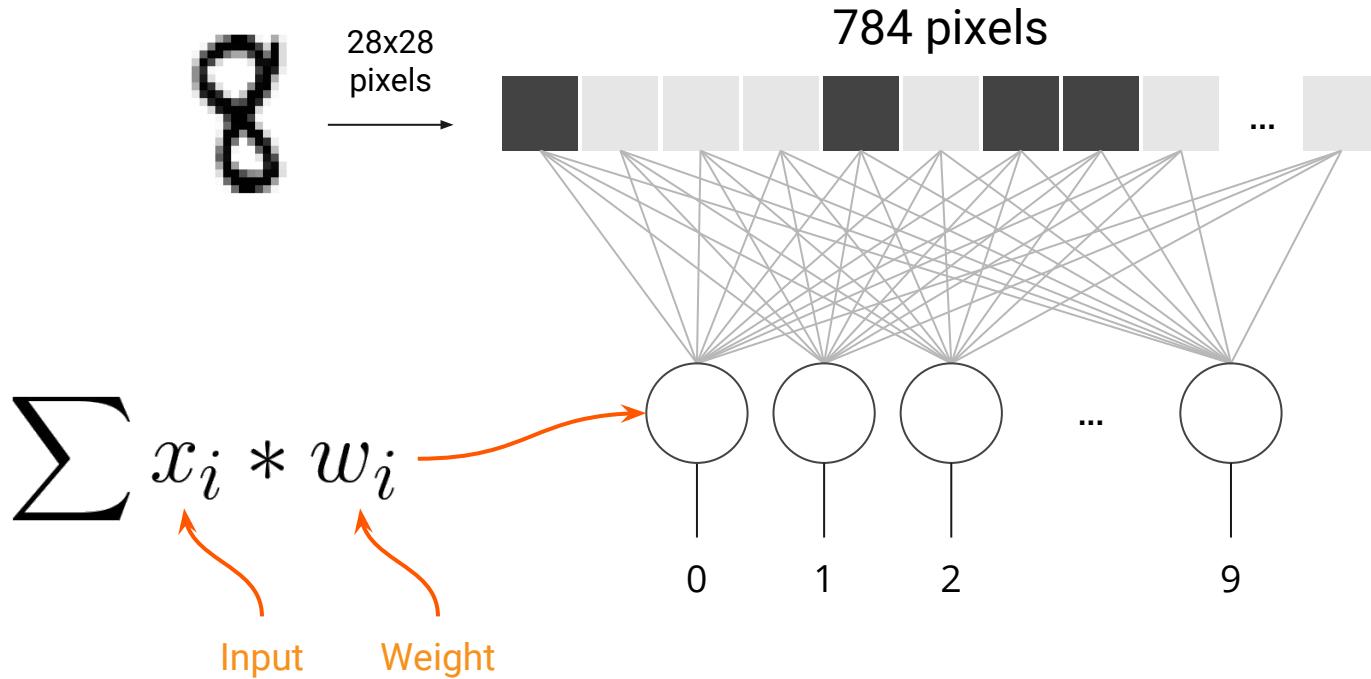
Test: 10,000

```
4 4 6 5  
2 3 6 1  
1 5 9 8  
1 2 9 7  
5 1 7 9  
4 6 2 2  
2 8 3 8  
0 0 3 0  
7 5 4 0  
7 1 3 0  
9 8 0 4  
8 5 7 8
```





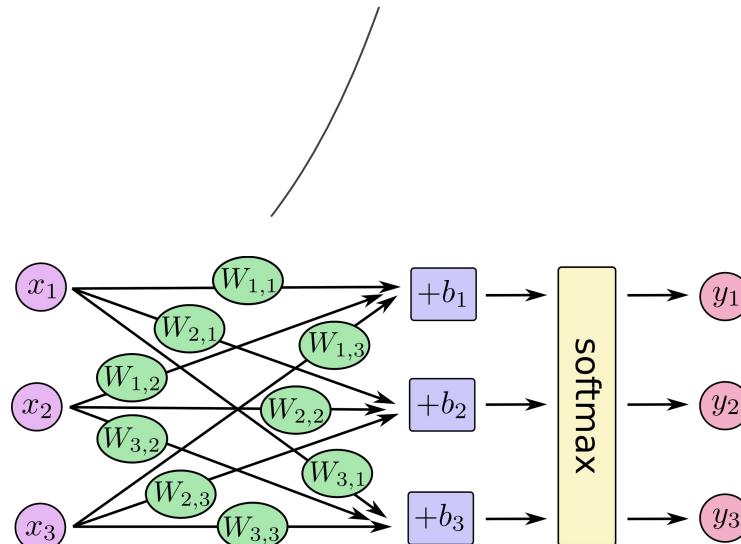
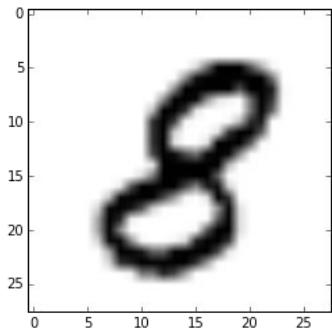




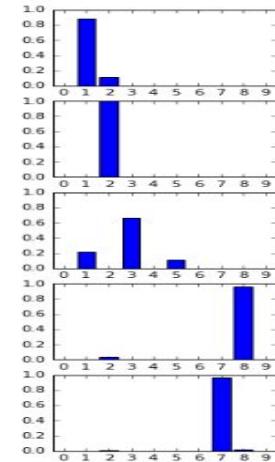
# Softmax

784 weights for each digit, so  $784 \times 10$  total connections

input vector  
(pixel data)



output vector  
(probability of  
each digit)



# Written as a matrix multiply

Weight matrix with dimensions [748,10]


$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix} = \text{softmax} \begin{pmatrix} W_{1,1}x_1 + W_{1,2}x_2 + W_{1,3}x_3 + b_1 \\ W_{2,1}x_1 + W_{2,2}x_2 + W_{2,3}x_3 + b_2 \\ W_{3,1}x_1 + W_{3,2}x_2 + W_{3,3}x_3 + b_3 \end{pmatrix}$$

# In TensorFlow

tensor shapes:  $X[batch\_size, 784]$      $W[784, 10]$      $b[10]$

$Y = \text{tf.nn.softmax}(\text{tf.matmul}(X, W) + b)$

matrix multiply

broadcast

# Placeholders

```
x = tf.placeholder(tf.float32, [None, NUM_PIXELS], name="pixels")  
y_ = tf.placeholder(tf.float32, [None, NUM_CLASSES], name="labels")
```

# Inference

```
w = tf.Variable(tf.zeros([NUM_PIXELS, NUM_CLASSES]), name="weights")  
b = tf.Variable(tf.zeros([NUM_CLASSES]), name="biases")  
y = tf.matmul(x, w) + b
```

# Loss and optimization

```
loss =  
    tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(y, y_))  
  
train_step =  
    tf.train.GradientDescentOptimizer(LEARNING_RATE).minimize(loss)
```

# Training

```
# Train the model
for i in range(TRAIN_STEPS):
    batch_xs, batch_ys = mnist.train.next_batch(BATCH_SIZE)
    sess.run(train_step, feed_dict={x: batch_xs, y_: batch_ys})
```

# Evaluation

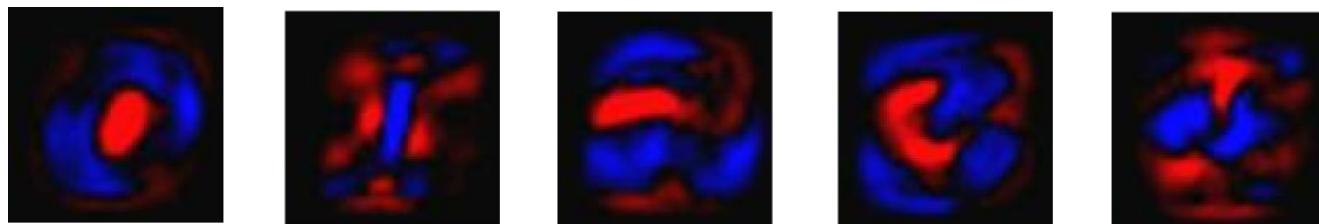
```
# Evaluate the trained model
correct_prediction = tf.equal(tf.argmax(y, 1), tf.argmax(y_, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))

print("Accuracy %f" % sess.run(accuracy, feed_dict={x:
mnist.test.images, y_: mnist.test.labels}))
```

**3\_basic\_mnist.ipynb**

**[goo.gl/nrdsxM](http://goo.gl/nrdsxM)**

# Visualizing the learned weights



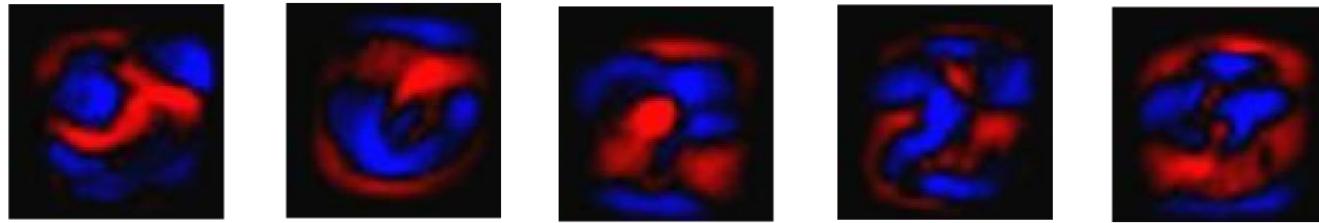
0

1

2

3

4



5

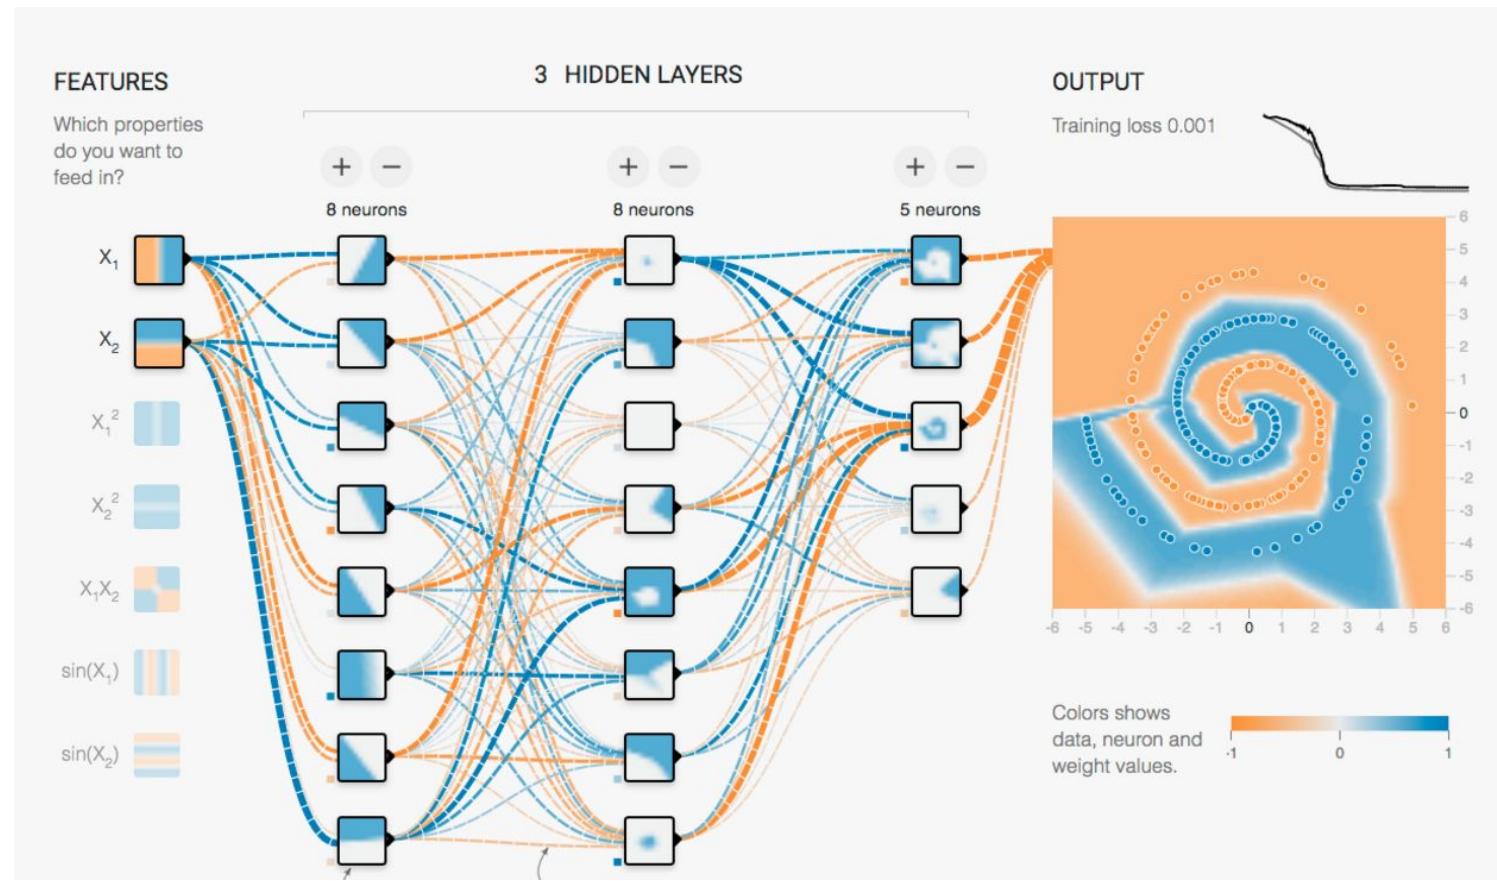
6

7

8

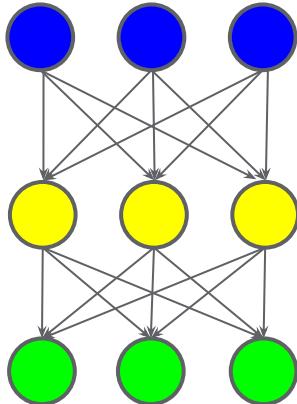
9

# Going Deeper



**More hidden layers = More hierarchies of features**

# Hidden layers



```
weights1 = weight_variable(NUM_PIXELS, HIDDEN1_UNITS)
```

```
biases1 = bias_variable(HIDDEN1_UNITS)
```

```
hidden1 = tf.nn.relu(tf.matmul(x, weights1) + biases1)
```

```
weights2 = weight_variable(HIDDEN1_UNITS, NUM_CLASSES)
```

```
biases2 = bias_variable(NUM_CLASSES, "biases2")
```

```
y = tf.matmul(hidden1, weights2) + biases2
```

# 4\_deep\_mnist.ipynb

[goo.gl/nrdsxM](http://goo.gl/nrdsxM)

# Progression and caveats

- Linear ~90%
- Two-layer DNN ~96%
- ConvNet ~99% (close to the state of the art on this toy problem)

## Should you focus on accuracy?

- What matters is designing a proper experiment.

1 <small><math>\times 1</math></small>	1 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	0	0
0 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	1 <small><math>\times 0</math></small>	1	0
0 <small><math>\times 1</math></small>	0 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved Feature

Convolution with  $3 \times 3$  Filter. Source:

[http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)

### Single depth slice

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

x

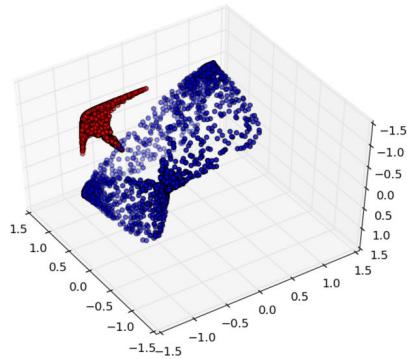
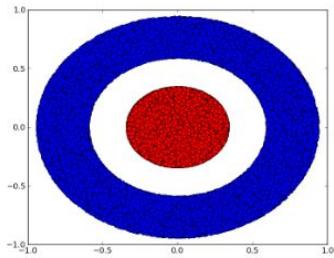
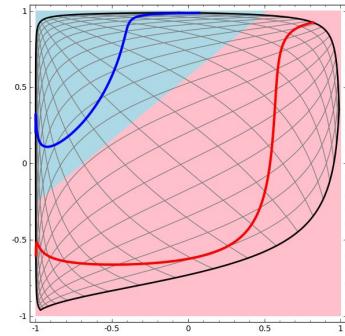
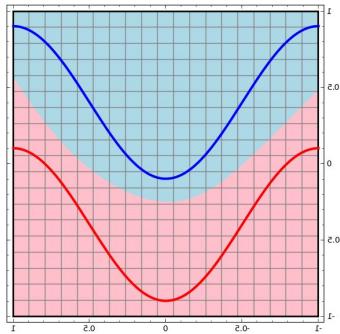
y

max pool with 2x2 filters  
and stride 2

6	8
3	4

Max pooling in CNN. Source: <http://cs231n.github.io/convolutional-networks/#pool>

Learn more: [colah.github.io](https://colah.github.io)



From: [Neural Networks, Manifolds, and Topology](https://colah.github.io), colah's blog

# Learn more: Stanford's cs231n

## CS231n Convolutional Neural Networks for Visual Recognition

These notes accompany the Stanford CS class [CS231n: Convolutional Neural Networks for Visual Recognition](#). For questions/concerns/bug reports regarding contact [Justin Johnson](#) regarding the assignments, or contact [Andrej Karpathy](#) regarding the course notes. You can also submit a pull request directly to our [git repo](#). We encourage the use of the [hypothes.is](#) extension to annotate comments and discuss these notes inline.

### Winter 2016 Assignments

[Assignment #1: Image Classification, kNN, SVM, Softmax, Neural Network](#)

[Assignment #2: Fully-Connected Nets, Batch Normalization, Dropout, Convolutional Nets](#)

[Assignment #3: Recurrent Neural Networks, Image Captioning, Image Gradients, DeepDream](#)

### Module 0: Preparation

[Python / Numpy Tutorial](#)

[IPython Notebook Tutorial](#)

[Terminal.com Tutorial](#)

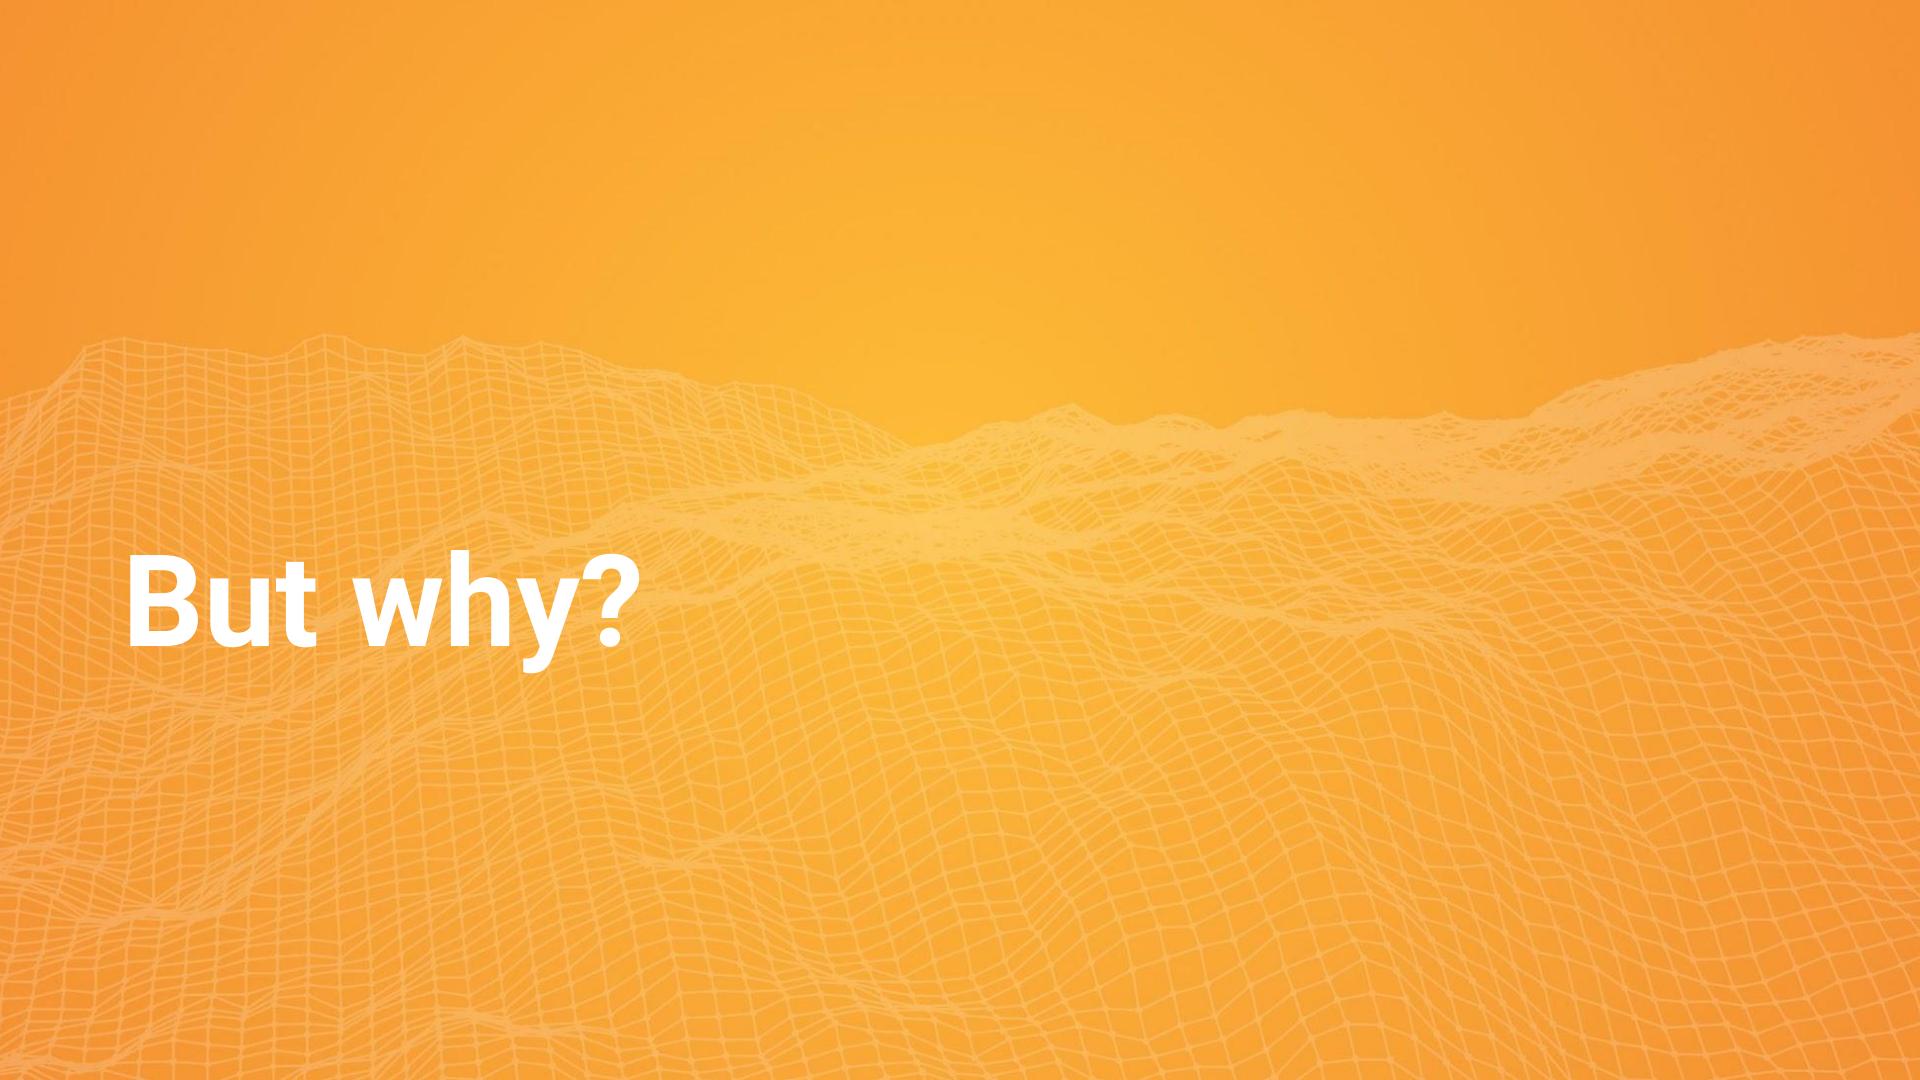
[AWS Tutorial](#)

### Module 1: Neural Networks

[Image Classification: Data-driven Approach, k-Nearest Neighbor, train/val/test splits](#)

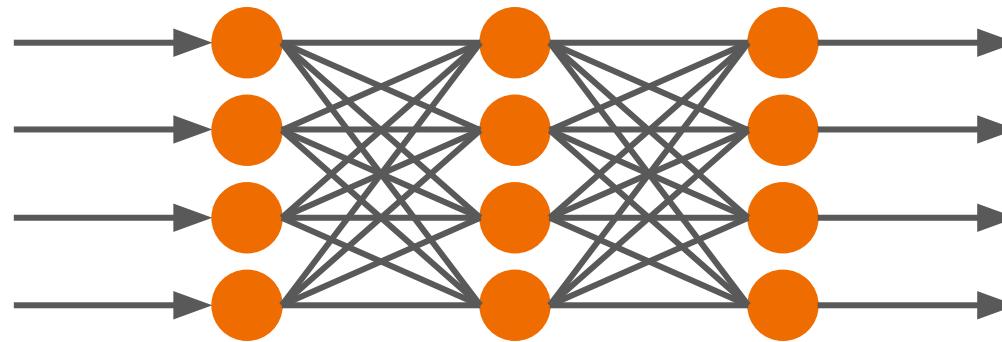
[L1/L2 distances, hyperparameter search, cross-validation](#)

[Linear classification: Support Vector Machine, Softmax](#)

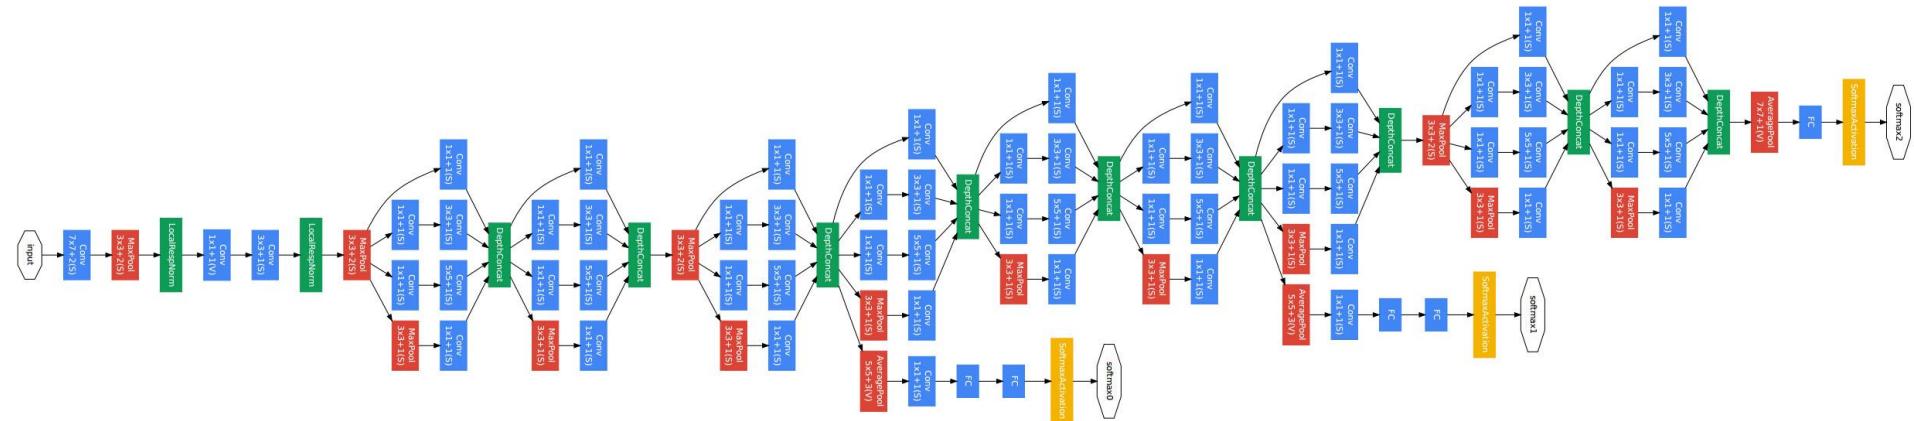
The background features a stylized landscape graphic composed of a grid of white lines on a solid yellow background. The grid forms a series of undulating hills and valleys, creating a sense of depth and perspective. The lines are thin and light-colored, allowing the yellow background to show through.

**But why?**

# Neural Networks, yesterday



# Remember Inception?



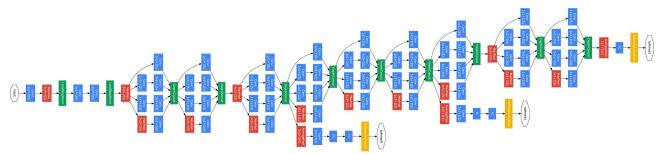
# Machine Learning gets complex quickly



Heterogenous  
systems



Distributed  
systems

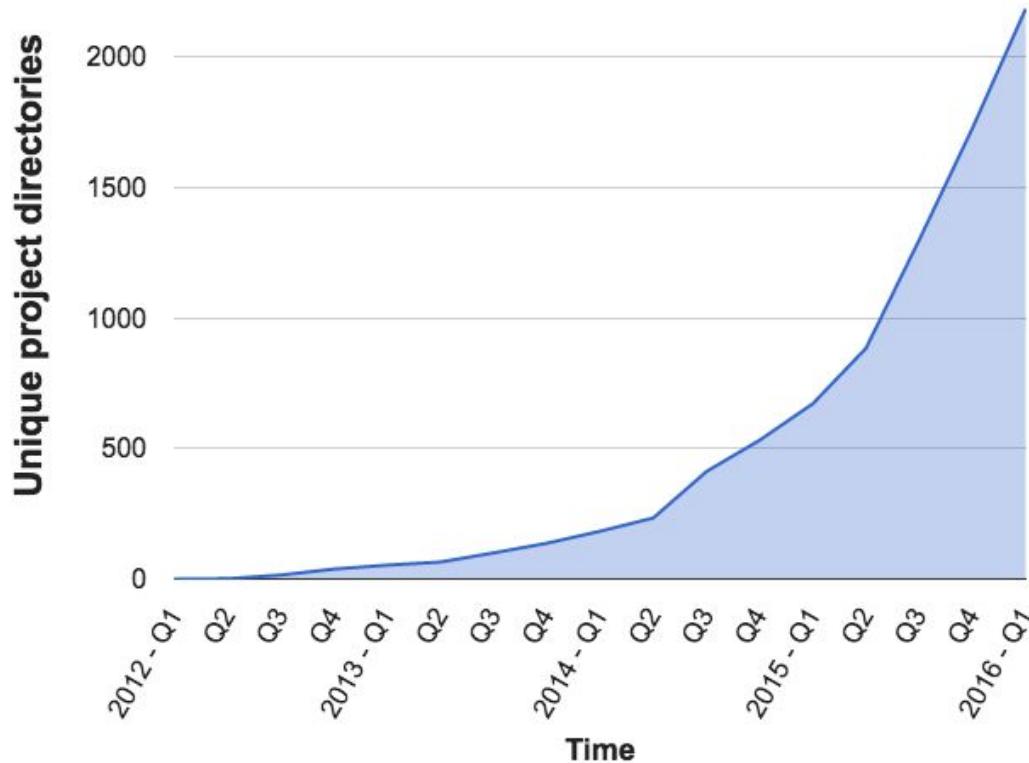


Modeling complexity

# Deep Learning at Google

# Growing Use of Deep Learning at Google

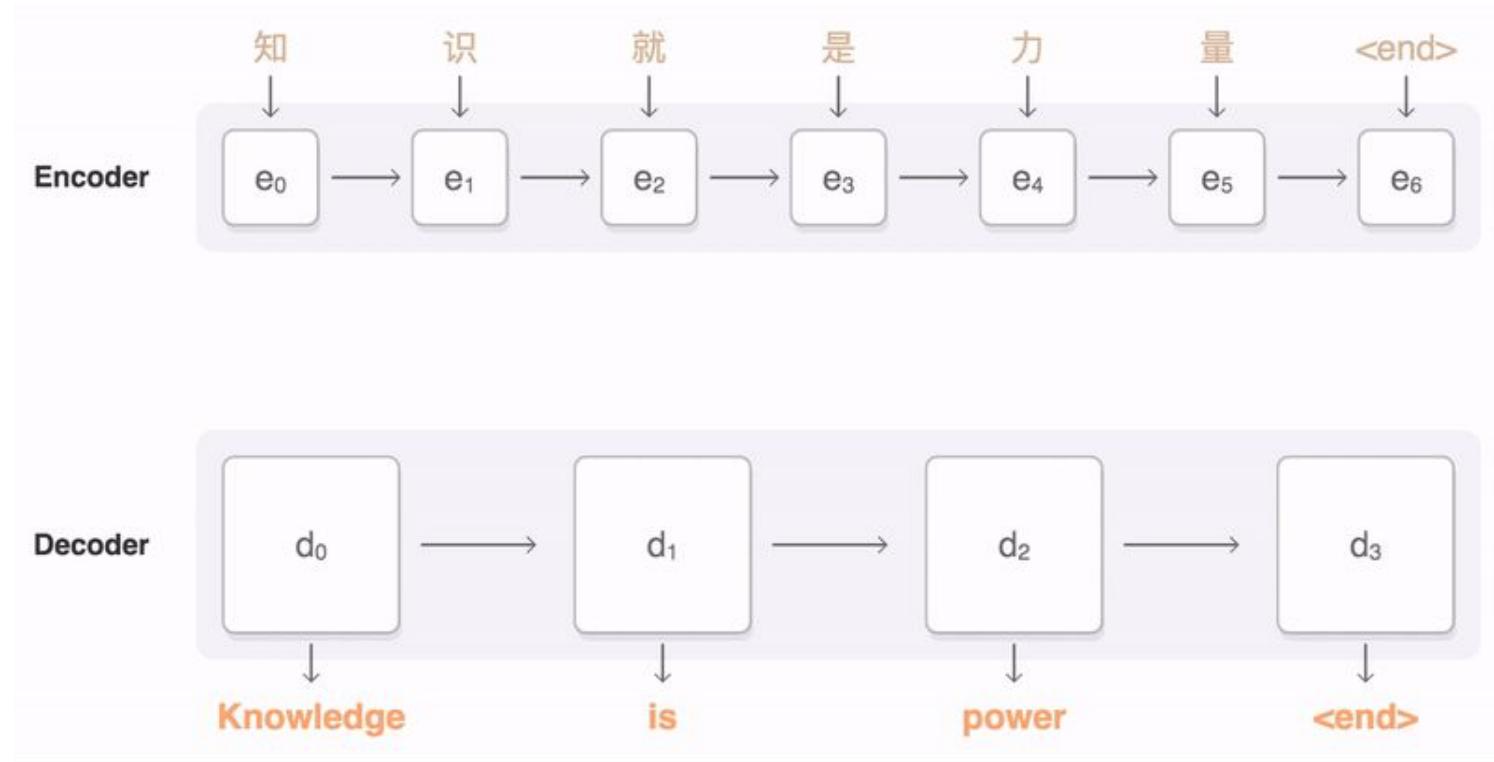
# of directories containing model description files

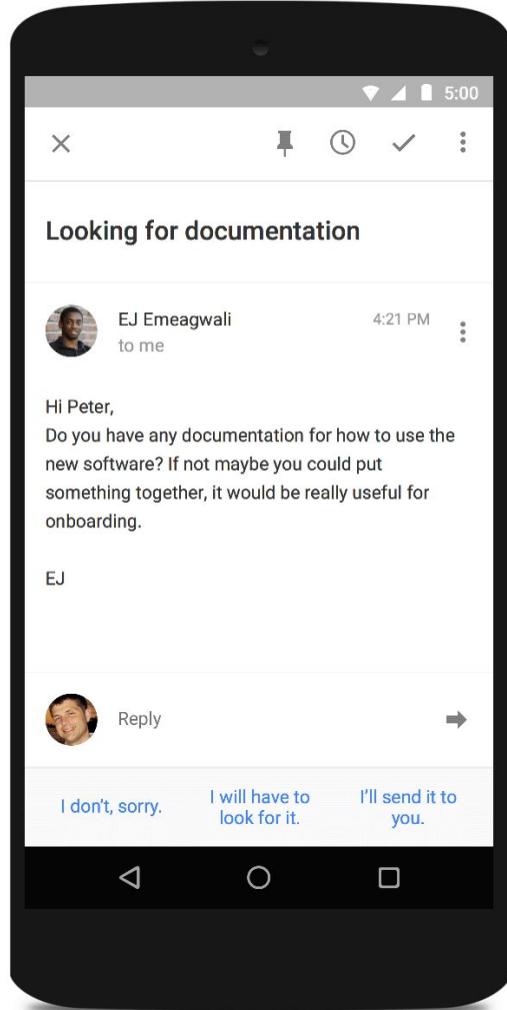




EXIT

# Google Neural Machine Translation



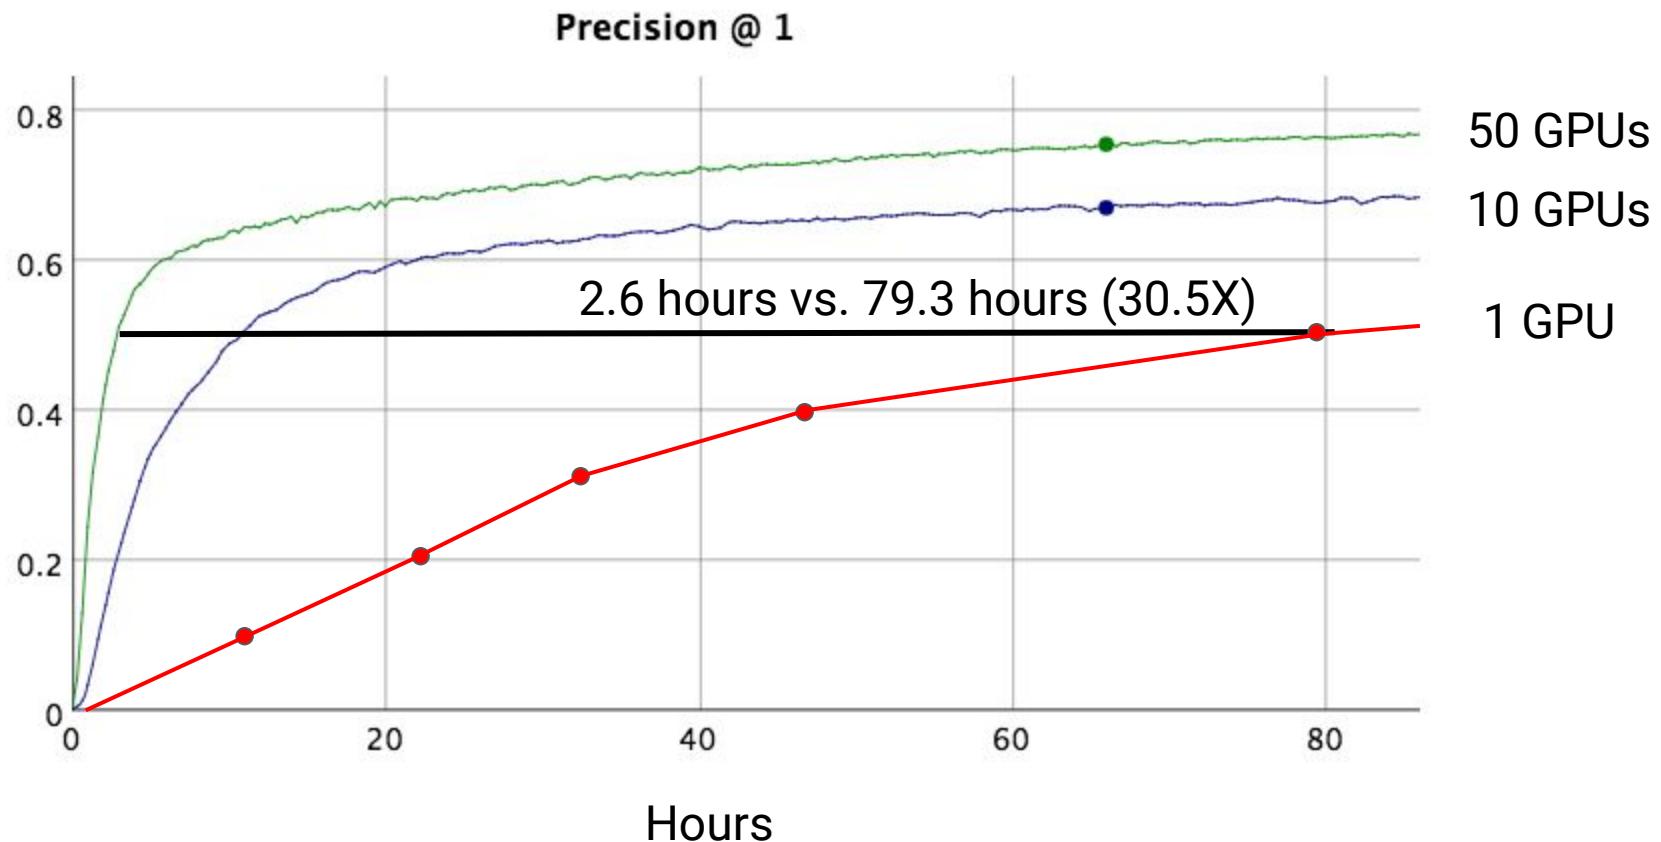


## Smart reply in Inbox by Gmail

10%

of all responses  
sent on mobile

# Image Model (Inception) Synchronous Training



# AlphaGo

BBC News Sport Weather iPlayer TV Radio More Search

Find local news

Home UK World Business Politics Tech Science Health Education Entertainment & Arts More

Technology

## Google achieves AI 'breakthrough' at Go

An artificial intelligence program developed by Google beats Europe's top player at the ancient Chinese game of Go, about a decade earlier than expected.

© 27 January 2016 | Technology

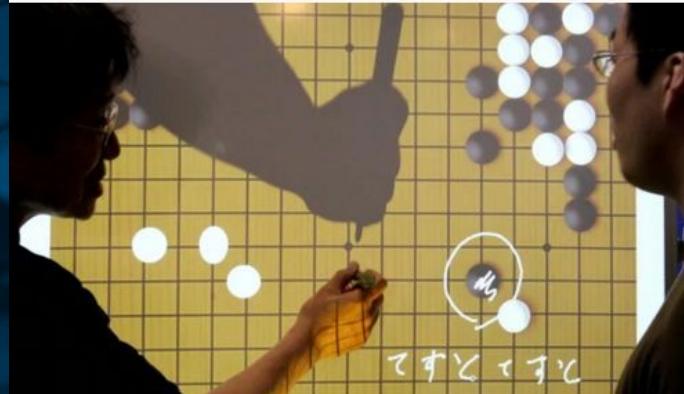
How did they do it?  
 What is the game Go?

Facebook trains AI to beat humans at Go



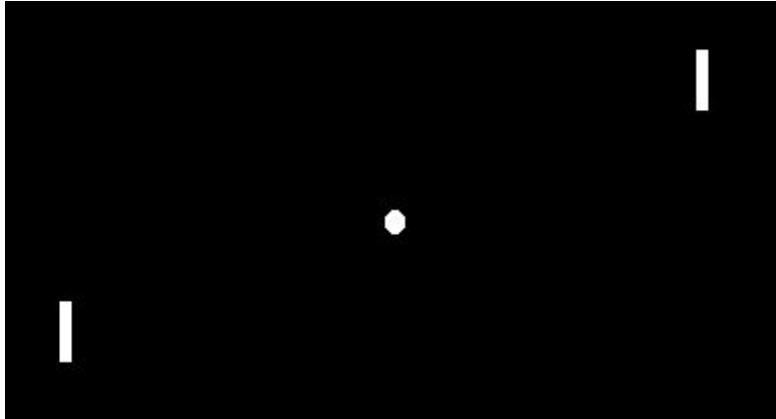
## Google's AI just cracked the game that supposedly no computer could beat

By Mike Murphy | January 27, 2016



Going up. (Reuters/Kiyoshi Ota)

Computers have slowly started to encroach on activities we previously believed only the brilliantly sophisticated human brain could handle. IBM's Deep Blue supercomputer beat Grand Master Garry Kasparov at chess in 1997, and in 2011 IBM's Watson beat former human winners at the quiz game *Jeopardy*. But the ancient board game Go has long been one of the major goals of artificial intelligence research. It's understood to be one of the most difficult games for computers to handle due to the sheer number of possible moves a player can make at any given point. Until now, that is.



# Next steps

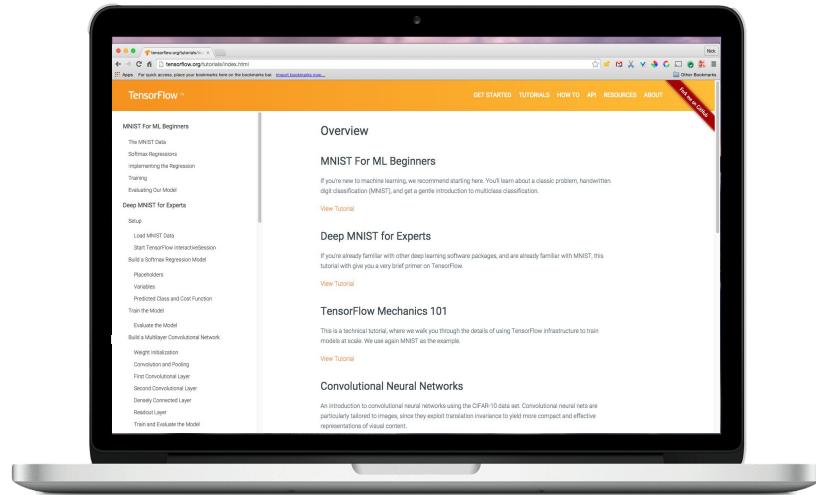
Tutorials and code  
[tensorflow.org](http://tensorflow.org)

Intro to Deep Learning with TensorFlow  
Udacity class [goo.gl/iHsslI](http://goo.gl/iHsslI)

Stanford's CS231n  
[cs231n.github.io](http://cs231n.github.io)

Udacity's Machine Learning Nanodegree  
[goo.gl/ODpXj4](http://goo.gl/ODpXj4)

Totally new to ML?  
Recipes [goo.gl/KewA03](http://goo.gl/KewA03)



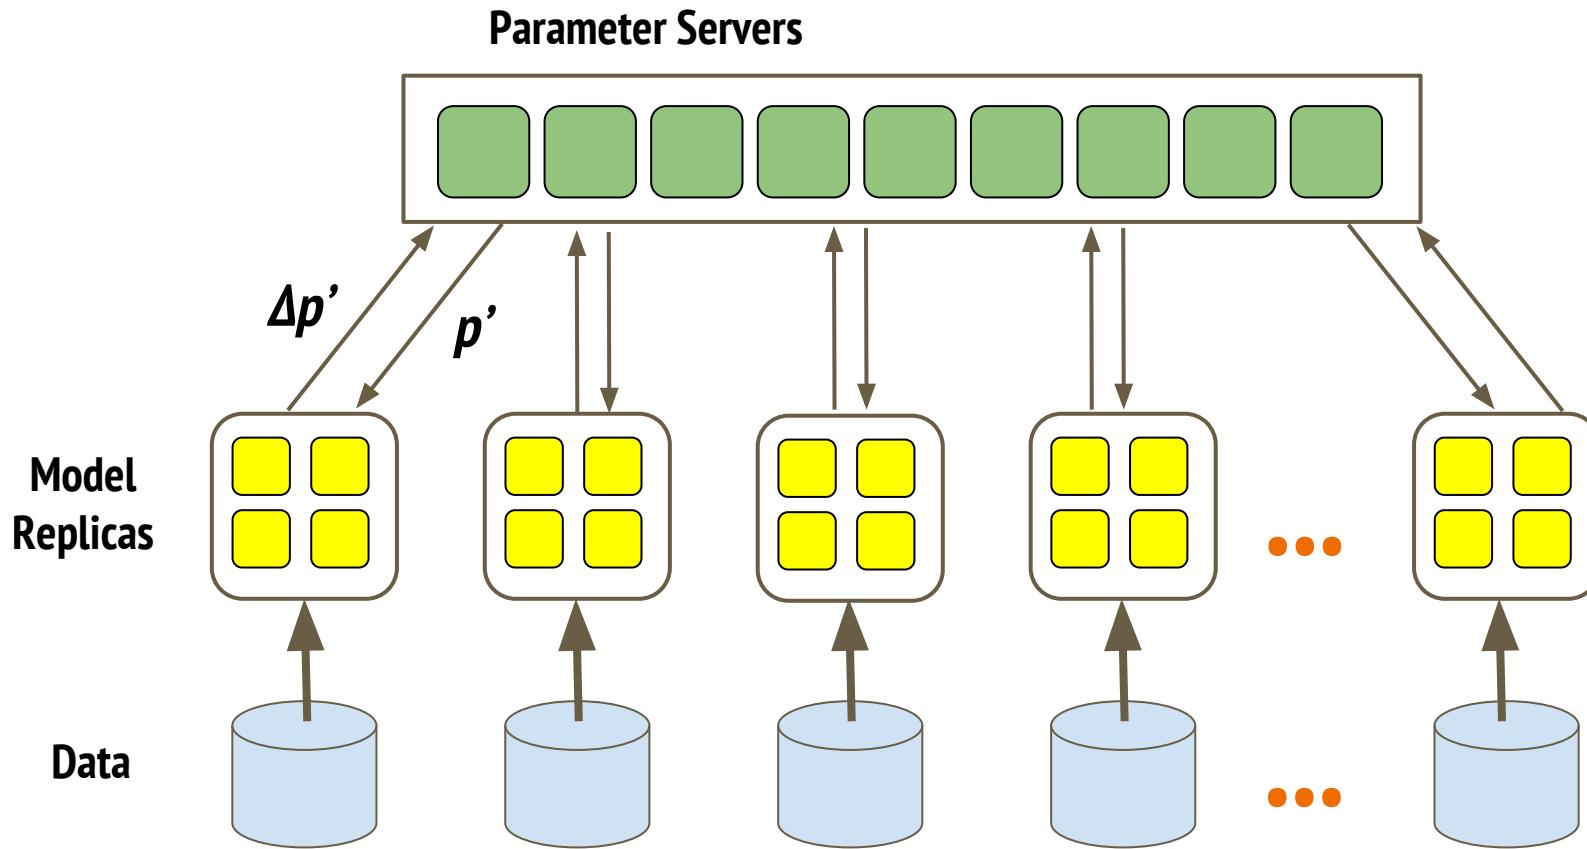
# Thank you and have fun!



Josh Gordon  
[@random\\_forests](https://twitter.com/random_forests)

# Extras

# Data Parallelism



# Input Pipelines with Queues

